

LLMs and IoT: A Comprehensive Survey on Large Language Models and the Internet of Things

Fatemeh Sarhaddi¹, Ngoc Thi Nguyen², Agustin Zuniga¹, Musfira Khan¹, Pan Hui³, Sasu Tarkoma¹, Huber Flores⁴, and Petteri Nurmi¹

¹University of Helsinki

²National University of Singapore

³Hong Kong University of Science and Technology (HKUST, GZ) and University of Helsinki

⁴University of Tartu

January 27, 2026

Abstract

The Internet of Things (IoT) has emerged as a technological cornerstone, enabling highly interconnected systems that integrate billions of diverse devices, with a combined data output estimated at hundreds of millions of terabytes per day. The rise of Large Language Models (LLMs) presents new opportunities for managing these devices, processing the data they generate, and addressing the technological and societal challenges posed by IoT ecosystems. However, the integration of LLMs into IoT also introduces significant challenges. This article offers the most comprehensive survey to date on the integration of LLMs into IoT ecosystems. We cover various technical layers, including device, software engineering, sensing, networking, data processing, privacy and security, and human interaction. We synthesize insights from over 300 articles, identify open research gaps, and highlight promising future directions, ranging from semantic reasoning and communication to the integration of quantum computing. By addressing the characteristics, challenges, and opportunities associated with LLM integration, this article serves as a foundational resource for researchers and practitioners aiming to enhance IoT functionalities with LLMs while navigating the associated complexities.

LLMs and IoT: A Comprehensive Survey on Large Language Models and the Internet of Things

FATEMEH SARHADDI^{*1}, NGOC THI NGUYEN^{*1}, AGUSTIN ZUNIGA^{*1}, MUSFIRA KHAN¹,
PAN HUI^{1,2}, SASU TARKOMA¹, HUBER FLORES³, AND PETTERI NURMI¹

¹Department of Computer Science, University of Helsinki, Finland

²Centre for Metaverse and Computational Creativity, Hong Kong University of Science and Technology (Guangzhou), Hong Kong

³Institute of Computer Science, University of Tartu, Estonia

^{*}These authors contributed equally to this work.

CORRESPONDING AUTHOR: Fatemeh Sarhaddi (e-mail: fatemeh.sarhaddi@helsinki.fi).

This research is supported by the Academy of Finland project 339614. The publication only reflects the views of the authors.

Preprint. This manuscript is under review at an IEEE Open Journal.

ABSTRACT The Internet of Things (IoT) has emerged as a technological cornerstone, enabling highly interconnected systems that integrate billions of diverse devices, with a combined data output estimated at hundreds of millions of terabytes per day. The rise of Large Language Models (LLMs) presents new opportunities for managing these devices, processing the data they generate, and addressing the technological and societal challenges posed by IoT ecosystems. However, the integration of LLMs into IoT also introduces significant challenges. This article offers the most comprehensive survey to date on the integration of LLMs into IoT ecosystems. We cover various technical layers, including device, software engineering, sensing, networking, data processing, privacy and security, and human interaction. We synthesize insights from over 300 articles, identify open research gaps, and highlight promising future directions, ranging from semantic reasoning and communication to the integration of quantum computing. By addressing the characteristics, challenges, and opportunities associated with LLM integration, this article serves as a foundational resource for researchers and practitioners aiming to enhance IoT functionalities with LLMs while navigating the associated complexities.

INDEX TERMS Large language models, LLMs, Internet of Things, IoT, artificial intelligence, AI, survey

I. INTRODUCTION

THE rapid proliferation of the Internet of Things (IoT) has led to the emergence of highly interconnected systems that integrate a wide range of devices. These systems provide significant benefits across various sectors, including environmental monitoring, healthcare, agriculture, and smart living environments [1]. Current estimates suggest that the number of connected IoT devices has surpassed 18 billion, with forecasts predicting 37.1 billion devices by 2032 [2]. Managing the vast network of IoT devices and their data is increasingly complex, especially in large-scale ecosystems. Conventional computing methods often struggle to cope with the diversity of IoT devices, the sheer volume of data, and the real-time demands of these environments. Consequently, there is a pressing need for advanced approaches capable

of efficiently managing massive volumes of heterogeneous IoT data while facilitating intelligent and context-aware interactions.

Large Language Models (LLMs), recognized for their advanced natural language processing capabilities [3], have emerged as powerful tools for facilitating natural and context-aware interactions. These sophisticated neural networks, pre-trained on extensive datasets, can understand and generate human-like language. Their proficiency in processing large datasets allows tasks such as text summarization, analysis, and generation. These capabilities open up opportunities for meaningful conversations, personalized recommendations, and enhanced decision-making [4].

The integration of LLMs has the potential to enhance IoT ecosystems by managing device diversity, processing

generated data, and addressing challenges within IoT environments. By leveraging LLM capabilities, IoT systems can achieve higher levels of intelligence and autonomy, leading to more effective operations. For instance, LLMs facilitate natural human-machine interactions [5], improve data interpretation through contextual understanding [6], [7], and automate tasks [8], [9]. Commercial IoT applications harnessing LLMs are also increasingly emerging, such as BMW integrating with Amazon's Alexa to provide natural language assistance during driving [10], and PepsiCo's Frito-Lay factories utilizing LLMs for predictive maintenance to identify possible faults and reduce downtimes [11]. Another example is Microsoft's Azure IoT platform, which supports LLM inference and deployment for industrial IoT applications [12].

However, the integration of LLMs presents both opportunities and significant challenges. These challenges include the resource limitations of IoT devices, privacy and security concerns related to sharing data with LLMs, and the unique characteristics and interactions among different devices in IoT environments. Recent advancements have begun to address some of these challenges through techniques such as model optimization [13], edge LLMs [14], instruction tuning [15], and prompting strategies [16]. While these developments enable the utilization of LLMs in a broader range of IoT scenarios, the overall research landscape regarding the effective integration of LLMs into IoT systems remains in its nascent stage and is insufficiently explored. There is a lack of information, among others, about specific methodologies for seamless integration, scalability, and the potential security and privacy implications of deploying LLMs across diverse IoT environments.

This article rigorously surveys the research landscape at the intersection of LLMs and IoT, covering key challenges, current uses across the different functional layers of the IoT ecosystem, future and emerging trends, and open gaps. Existing surveys focus primarily on optimizing LLMs or integrating them into specific IoT components [17]–[22], without providing a holistic overview of the entire IoT ecosystem. IoT applications are characterized by complex interactions across different functionalities. Without considering a holistic view of the entire ecosystem, the full potential of LLMs as well as many challenges and gaps may be overlooked. Our survey addresses this limitation in previous research by analyzing the challenges and opportunities associated with LLM integration into IoT systems. We review over 300 articles to synthesize advancements in LLM integration across IoT layers, from devices and software engineering to sensing, networking, and data processing. We also discuss security, privacy, and human interaction in LLM-powered IoT systems and their implications. Finally, we highlight future directions and emerging trends for research, as well as open research gaps identified from our findings. Addressing these gaps is crucial for maximizing LLM integration's potential and enhancing IoT capabilities. This

survey serves as a foundational resource for researchers and practitioners looking to improve IoT functionalities with LLMs while navigating associated challenges.

II. SCOPE AND METHODOLOGY

Research on Large Language Models (LLMs) has surged exponentially in recent years. A Google Scholar search for "large language models" reveals 149 000 publications from 2023 to January 9, 2026, covering a wide range of topics, from foundational model architectures (e.g., GPT models [23], [24], LaMDA [25], LLaMa [26], BERT [27], and others) to innovative use cases across various application domains (e.g., healthcare [28]–[30], education [31], [32], environmental monitoring [33], [34], smart homes and cities [35], [36], and industrial IoT [37], [38]).

When narrowing the query to focus specifically on articles related to IoT, the volume of publications is smaller but shows a notably growing trend. Specifically, the query "Large Language Models" "LLMs" "Internet of Things" "IoT" results in 8500 articles over the same period, with 530 results in 2023, 3000 in 2024, 6480 in 2025, and 171 as of January 9, 2026. Most of these studies focus on the research challenges of integrating LLMs with IoT technologies [39], or they provide specific examples of LLM applications within IoT contexts [36], [40]. However, a comprehensive understanding of the potential that LLMs can offer and the challenges they pose for IoT systems is notably lacking. Indeed, the integration of LLMs into IoT presents unique challenges and limitations, such as memory, energy, and other resource constraints, hardware heterogeneity, and privacy and security challenges. These factors complicate the deployment and use of LLMs as part of IoT ecosystems.

Our survey addresses this gap by providing a comprehensive and systematic review of LLMs across the entire IoT architecture, conducting a rigorous analysis of the open research challenges in the field, and providing insights into how LLMs can be tailored to meet the specific demands of IoT systems. We emphasize both the breadth (covering all IoT layers) and depth (examining cross-layer interactions and challenges) of LLM integration, which distinguishes our work from prior surveys. This allows us to identify challenges and dependencies emerging from the integration of LLMs across all IoT layers, offering a holistic understanding of optimal deployments, resource allocation, and security and privacy considerations. Overall, our work bridges the existing knowledge gap and provides a foundational understanding of how LLMs can be effectively integrated into IoT frameworks.

A. RELATED SURVEYS

Table 1 provides a structured comparison of the most relevant surveys across several dimensions: LLM features, application domains, and coverage of the main IoT functional layers. Surveys focusing on *LLM features* delve into specific char-

TABLE 1: Comparison between related surveys and this survey, focusing on LLM background, domain, and IoT layers.

Reference	Year	Article corpus	LLM background				Domain				Devices		Software Engineering		Sensing		Networking			Data Processing			Privacy & Security		Human Interaction		Challenges / Limitations	Opportunities	
			Architecture	Prompt engineering	Fine-tuning	Reasoning	General	LLMs and IoT	Applications	Code generation	Task planning / management	On-device	Distributed	Development	Automation	Data generation	Context-aware sensing	Networking Technologies	Device Communication	Offloading	Analysis	Evaluation	Management	Privacy	Security	Intelligent UI			Personalization
Bai et al. [17]	2024	✓	✓	✓	✓	✓			✓	✓								✓								✓	✓		
Chang et al. [41]	2024	✓		✓		✓	✓		✓					✓						✓						✓	✓		
Cui et al. [42]	2024	✓		✓	✓	✓		✓														✓	✓			✓	✓		
Das et al. [43]	2025	✓		✓	✓		✓		✓	✓													✓	✓			✓	✓	
Ferrara et al. [21]	2024	✓		✓	✓			✓	✓					✓	✓				✓				✓	✓			✓	✓	
Friha et al. [44]	2024	✓		✓		✓	✓		✓		✓	✓					✓		✓				✓	✓			✓	✓	
Guo et al. [45]	2024	✓			✓				✓					✓					✓				✓	✓			✓	✓	
Hu et al. [46]	2024		✓		✓			✓		✓	✓								✓						✓	✓			
Jin et al. [47]	2023	✓		✓			✓		✓				✓						✓						✓	✓		✓	
Kök et al. [39]	2024				✓			✓		✓									✓				✓	✓			✓	✓	
Lawrence et al. [48]	2024			✓					✓					✓					✓				✓	✓			✓	✓	
Li et al. [19]	2024	✓			✓			✓								✓			✓				✓	✓	✓	✓	✓	✓	
Minaee et al. [4]	2024		✓		✓		✓			✓																	✓	✓	
Naveed et al. [18]	2023			✓	✓				✓	✓		✓	✓		✓				✓	✓							✓	✓	
Qu et al. [49]	2025	✓	✓	✓	✓		✓		✓		✓	✓						✓	✓				✓	✓			✓	✓	
Wang et al. [50]	2024	✓		✓			✓																✓	✓			✓		
Yao et al. [3]	2024	✓		✓	✓				✓	✓		✓	✓					✓					✓	✓			✓	✓	
Zhang et al. [15]	2023	✓	✓		✓		✓		✓							✓													
Zhang et al. [22]	2024	✓	✓	✓				✓		✓				✓						✓							✓	✓	
Zhou et al. [51]	2024	✓	✓	✓	✓		✓		✓		✓						✓		✓					✓					
Zhu et al. [52]	2023	✓	✓	✓			✓												✓	✓							✓	✓	
Our survey	2025	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	
Total of Surveys		15	12	19	18	5	16	6	13	11	5	7	5	3	2	7	4	3	6	5	10	5	2	13	13	3	5	19	16

acteristics such as resource efficiency, architecture, training strategies, and prompting techniques. For instance, some reviews explore various prompting techniques for natural language processing (NLP) [16]. The use of LLMs has been extensively documented across numerous application domains, including healthcare [21], education [41], mental health [45], [48], personal assistance [19], and gaming [46]. However, these surveys address domain-specific issues and overlook IoT-specific challenges that arise from their integration into IoT applications.

In terms of IoT architecture, surveys on *IoT Devices* and *IoT Software Engineering* address the development of LLM-based agents and instruction tuning for GPT models, while those focused on *IoT Sensing* investigate LLMs for sensor data retrieval and model tuning. However, there is a notable scarcity of literature on *IoT Networking*, with current surveys primarily covering telecommunication networks or edge computing [39], [44], [49], [51]. Surveys on *IoT Data Processing* evaluate methods for analyzing sensor data and multimodal datasets using LLMs, and those addressing *IoT Privacy and Security* examine vulnerabilities associated with LLMs and their implications for privacy and security at both the system and user levels. While these surveys have explored specific LLM characteristics or isolated IoT layers, they overlook the significant interactions across the entire IoT landscape. A comprehensive understanding of these interactions is essential for effectively harnessing LLMs in IoT applications, as many fundamental architectural layers and key components—such as hardware and software, sensing and data analysis, and communications and networking—remain largely unexplored.

As this comparison indicates, existing surveys primarily address specific challenges or application domains within specific IoT domains or layers, lacking a holistic view of the overall IoT ecosystem. This significantly restricts the insights they can derive about interactions between different IoT components and subsystems. Our survey addresses this critical gap by systematically investigating the potential of LLMs across all IoT layers and identifying open challenges both within and between these layers. By offering a holistic perspective on the integration of LLMs into various IoT layers, this paper provides a comprehensive understanding of their applications, potential benefits, limitations, and the challenges associated with utilizing LLMs in the IoT landscape.

B. SCOPE OF THE SURVEY

This survey provides a comprehensive understanding of the current state-of-the-art regarding the various roles of LLMs within IoT systems and to identify critical research gaps. Specifically, we explore the potential benefits of LLMs for IoT systems, investigate how these systems can effectively leverage LLM capabilities to enhance their development processes, and assess the existing challenges, limitations, and potential risks associated with integrating LLMs into IoT environments. We start with a brief overview of current LLM technologies in Section III and briefly cover the challenges associated with the deployment of LLMs in Section IV. The core of the survey, presented in Section V, is structured to cover the primary technological layers of IoT, including devices, software engineering, sensing, networking, data processing, privacy and security, and human interaction. An overview of the overall scope of the survey is shown in




FIGURE 1: LLMs and Their Interaction with Functional Layers of IoT Applications. Sensing, networking, and data processing are classified as specialized functional layers, each addressing distinct aspects of IoT operations. In contrast, devices, software engineering, privacy and security, and human interaction are cross-cutting concerns that influence the entire IoT ecosystem, affecting all systems and deployments. Our survey scope addresses LLMs integration within all functional layers.

Figure 1. Section VI synthesizes our analysis, presenting future directions and emerging trends as well as identifying open gaps for future research. Finally, Section VII concludes the article.

C. METHODOLOGY

The article selection and analysis process for this survey was conducted in three stages: article identification, preliminary screening, and comprehensive article analysis. Each stage was designed to systematically identify the most relevant and best quality articles to include in the final survey.

Stage 1: Article Identification. In this first stage, we conducted a systematic search across three major academic databases: ACM Digital Library, IEEE Xplore, and Google Scholar. This search was first performed in September 2024, and was repeated on 15 February 2025. We employed the keywords and their possible combinations of *Large Language Models*, *LLMs*, *Internet of Things* and *IoT*. This approach ensures that we captured articles discussing LLMs in the context of IoT from different perspectives. We supplemented the results by searching for follow-up works and articles that cite and were cited by them. Duplicate versions

of the same article (e.g., preprints, published versions, or conference presentations) were removed.

Stage 2: Preliminary Screening. In the second stage, we evaluated the articles identified in the first stage. The preliminary screening involves reviewing abstracts to assess their relevance on several key aspects: (i) LLM architectures: technical details, datasets used for model training, interfaces, and (ii) the utilization of LLMs in IoT: knowledge acquisition, feature engineering, inference, recommendation, privacy, security, and user interaction. Each article was then tagged as *potentially relevant* or *irrelevant* based on its alignment with the survey’s scope and these key aspects. As such, we excluded articles that merely mention LLMs in passing or as part of broader discussions without a focused analysis. Additionally, articles that focus solely on theoretical aspects of LLMs or IoT without practical applications were discarded. This criterion was essential to keep our survey focused on research providing insights and solutions to real-world challenges. To finalize the selection of articles, we employed a majority decision-making process among the authors. This collaborative approach ensures that the final selection reflects a consensus on the relevance and quality of the articles, allowing us to compile a comprehensive and

insightful survey. Our final corpus includes over 300 articles. Links to these articles, organized by the functional IoT layer they address, are available in a separate GitHub repository¹.

Stage 3: Comprehensive Article Analysis. In the final stage, we conducted a comprehensive analysis of the articles in the corpus. We meticulously reviewed each article to identify and label key concepts (e.g., feature extraction, preprocessing time-series data) and phrases (e.g., generating IoT application scenarios) relevant to our survey scope. This step allows us to capture a wide range of information without imposing predefined categories, fostering a broader understanding of the articles on how IoT systems can leverage the capabilities of LLMs. We then organized the extracted information into categories and explored the relationships among them to uncover insights on the complexities and nuances of LLM integration with IoT systems. By systematically examining both the breadth and depth of the literature, our analysis not only provides a comprehensive understanding of the current research landscape surrounding LLMs and IoT but also identifies key challenges and gaps within the existing literature, guiding future investigations and practical applications in the integration of LLMs and IoT. Lastly, building on this foundational body of literature, we identified key ideas, research directions, and emerging technologies that can further advance the integration of LLMs and IoT. This broader perspective highlights how their convergence is driving innovation across multiple fields. Taken together, our survey serves as a foundational resource for researchers and practitioners in the fields, equipping them with essential insights to drive innovation for addressing existing challenges and explore new avenues for effective integration of LLMs and IoT.

III. LARGE LANGUAGE MODELS

Large Language Models (LLMs) are advanced neural networks that have been pre-trained on extensive datasets containing diverse forms of data, including text, code, and more. This pre-training equips them with the ability to (i) deeply understand and extract information from human-like text, (ii) generate coherent and contextually relevant text, (iii) infer context and demonstrate contextual awareness, and (iv) provide recommendations or instructions [3]. Notable examples of LLMs include the Generative Pre-trained Transformer (GPT) series, BERT [27], XLNet [53], ChatGLM [54], LLaMA [26], PaLM [55], Claude [56], DeepSeek [57], Gemini [58], and Grok-3 [59]. We refer to [60] for an up-to-date list of LLM tools and their performance in different tasks. There have also been efforts to create specialized models that address specific aspects of IoT or target IoT-specific industries. For example, Google SensorLM is a research project focusing on creating a foundation model for wearable data [61]. Examples of enterprise LLMs, in turn,

include SymphonyAI [62], H2O.ai [63], [64], Marble [65], and research initiatives by Microsoft to establish industry-specific foundation models [66].

The development of LLMs is rooted in the evolution of language modeling (LM) techniques. Initially, language modeling served as a foundational approach for understanding and manipulating human language within the field of natural language processing (NLP). Early models utilized statistical methods to calculate the probability distributions of word sequences [67]. Over time, these models evolved into statistical language models that predict the next word in a sequence. With the advent of deep learning, SLMs transitioned into neural language models (NLMs), which emphasize the interaction between human and computer language [68]. Pre-trained language models (PLMs) represent a specific category of language models that leverage large datasets to transfer knowledge from previously learned tasks to new applications [69]. PLMs with substantial amounts of data and a large number of parameters (often exceeding 6 billion) are commonly referred to as Large Language Models (LLMs) [68]. This evolution not only improved the performance of language models but also broadened their applicability across various domains, marking a significant milestone in the field of NLP. More recently, Large Action Models (LAMs) are special variants of LLMs that extend LLM capabilities to deliver agents that autonomously perform real-world actions [70].

At the core of LLMs lies the transformer neural network architecture, which operates on *tokens*. A token is a unit of input data that can be transformed into a vector representation, serving as the processing unit within the transformer model. The process of converting input data into tokens is referred to as *tokenization*, where the tokens can represent characters, words, or other elements. This tokenization enables efficient processing and analysis of text. The transformer architecture employs two key mechanisms that facilitate the learning of dependencies and relationships between tokens. The first mechanism is *self-attention*, which allows the model to weigh the importance of each token relative to others, capturing internal correlations and contextual relationships. The second mechanism, *multi-head attention*, enables the model to perform multiple self-attention operations in parallel, with each "head" focusing on different aspects of the input. By capturing a broader range of dependencies, multi-head attention enhances the model's ability to understand complex relationships within the data [71].

The transformer model itself consists of two main components: the *encoder* and the *decoder*. The encoder processes the input data, extracts features, and learns the relationships between tokens using self-attention mechanisms. It embeds contextual information along with the input sequence into a series of vectors for the decoder to utilize. The decoder generates output data sequentially based on the input received from the encoder. In addition to multi-head self-attention and

¹The repository is made public during the camera-ready phase. Screenshots of the repository are included in the cover letter.

feed-forward networks, the decoder incorporates encoder-decoder attention, which connects it with the encoder. This connection enables the decoder to compute attention scores for all tokens in the encoder's output, determining which parts of the input sequence are relevant for generating the output.

During a query, each token is transformed into three vectors: a Query vector, a Key vector, and a Value vector. The Query vector contains information of interest, the Key vectors represent each token's unique characteristics, and the Value vector holds the actual information associated with each token. By comparing the Query vectors with the Key vectors, the model computes attention scores that indicate the relevance of each token in relation to others. Throughout this process, self-attention is iteratively applied across layers, capturing dependencies that are stored in memory as attention key-value tensors, collectively referred to as the key-value (KV) cache [72]. The storage and retrieval of the KV cache and model parameters from memory scale the inference cost quadratically with input sequence length [73]. In fact, the KV cache for a single request can be as large as 1.6GB [74]. Insufficient management of the KV cache may exhaust memory that could otherwise be allocated for processing incoming requests.

To mitigate the memory footprint associated with LLMs, the KV cache can be offloaded to CPU memory, either on the same device or to another device over the network [75]. Additionally, previously computed key/value vectors can be reused for generating the current token through techniques such as KV cache compression [76]. However, offloading the KV cache, whether locally or across a network, can introduce latency, which may hinder performance. This challenge is especially pronounced in resource-constrained environments, such as Internet of Things (IoT) devices, where limited CPU memory complicates the management of large caches effectively. Consequently, there is a pressing need for innovative memory management strategies to ensure efficient deployment of LLMs in such settings.

Traditional LLMs utilize both an encoder and a decoder to process input and generate output. However, the demands for efficiency, simplicity, and task specialization have led to the rise of encoder-only and decoder-only models, each optimized for specific applications. Encoder-only LLMs, such as BERT [27] and ALBERT [77], utilize the bidirectional self-attention mechanism of the encoder component of the transformer architecture. This allows them to capture contextual information and dependencies from the input data effectively, resulting in rich contextual embeddings that represent the semantics of the input. Encoder-only models excel in tasks that require a deep comprehension of the input data, making them efficient for downstream tasks such as sentiment analysis, feature extraction, and classification [78]. In contrast, decoder-only LLMs, such as GPT [24], Claude [56], and LLaMA [26], utilize an auto-regressive self-attention mechanism within the decoder component of the transformer

architecture. This design enables decoder-only models to generate text sequentially based on previously predicted tokens to maintain coherence and context. This makes decoder-only models particularly well-suited for tasks that involve creating coherent and contextually relevant sequences of text, such as conversational agents, contextual reasoning, creative writing, code generation, and summarization [78]. These architectural variations offer flexibility to tailor LLMs to specific task requirements. Tasks focusing on understanding and analyzing data without the need to generate responses are best served by encoder-only models, whereas tasks that necessitate real-time response generation without the need for understanding or analyzing the data benefit from decoder-only models. Accordingly, tasks that require both capabilities should utilize models that integrate both the encoder and decoder components.

IV. CHALLENGES IN LLM AND IOT INTEGRATION

The integration of LLMs with IoT systems poses significant challenges due to the complexities of both technologies. We next discuss the key challenges in the use of LLMs, applicable across the entire IoT ecosystem. Issues related to individual IoT layers are discussed in the next section. While these challenges are not exclusive to IoT, understanding them is crucial for gaining insights into the research landscape and ensuring the effective integration of LLM and IoT technologies.

Resource Constraints and Sustainability. LLMs require immense computational power, memory, and energy, while many IoT devices have limited resources. This can lead to inference latency, high power consumption, and challenges in storing LLM parameters. Techniques to optimize efficiency of LLMs include model optimization techniques discussed below, and hardware acceleration discussed in Section V-A. Besides the model itself, the key-value cache used by LLMs also demands significant memory, exacerbating resource demands. Strategies to mitigate memory usage include efficient compression techniques and strategic cache management [79], [80]. Many IoT devices are battery-powered, and the high energy requirements of LLMs pose sustainability concerns [81]. Optimizing resource utilization is critical for operational efficiency and sustainability in LLM applications within IoT environments.

Model Optimization. Model optimization methods focus on reducing LLM model size and resource demands, but often do so at the cost of accuracy. Key methods include pruning, quantization, and knowledge distillation [82], [83]. Pruning eliminates less significant parameters, layers, or channels, making it well-suited for edge and on-device deployments [84]–[86]. IoT-specific pruning methods, such as DeepIoT [87] and FastDeepIoT [88], focus specifically on achieving better compression-accuracy trade-offs while ensuring compatibility with resource-constrained devices. Quantization reduces resource demands by adjusting parameter precision, which is particularly effective on low-

end devices [13], [89], [90]. Modern quantization often employs mixed-precision modeling, utilizing 16-bit and 32-bit representations for different parts of the model [91]. While this accelerates inference and training, it may lead to instability, resulting in loss spikes due to inaccuracies in gradient estimates. Knowledge distillation transfers knowledge from larger models to smaller ones [92], though it can reduce model diversity and adaptability [93], [94]. The DeepSeek model family exemplifies how these techniques can lower inference requirements for conventional hardware [57], [95]. Another option for optimizing model efficiency is to use neural architecture search (NAS) techniques to identify a smaller model architecture that closely resembles the original while meeting the resource constraints of the target platform, with promising results reported in applying NAS to LLMs [96]. However, a major drawback of these techniques is that finding the optimal model architecture often requires significant computational resources, which can limit the ability to adapt models on demand.

Prompt Engineering. Prompt engineering involves crafting rules and instructions to maximize LLM capabilities based on specific data types. Common methods include zero-prompting, which simplifies interactions by using an LLM without a detailed prompt, and zero-shot learning, which enables models to predict unseen classes using pre-learned attributes [97], [98]. This approach has been effectively applied in IoT domains, such as healthcare and smart homes, by incorporating contextual information into prompts [99]–[102]. Generalized zero-shot prompting enhances adaptability by capturing semantic relationships, while few-shot learning improves performance by leveraging a small number of labeled examples, often outperforming zero-shot prompting [100], [103]. Advanced techniques, such as Chain of Thought (CoT) prompting, provide step-by-step explanations to improve reasoning quality [104]. CoT has been successfully applied in activity recognition and task planning in smart homes [36], [105], while self-consistency (SC) prompting enhances logical consistency, especially in complex fields like healthcare [100]. Other methods, like Program of Thought (PoT) prompting, enable LLMs to generate code for complex tasks [106]. Recently, token-budget-aware methods aim to compress LLMs' reasoning process by considering the token budget for each prompt, for example, Han et al. [107] predict an adjustable token budget for each CoT prompt, SelfBudgeter [108] utilizes reinforcement learning to predict the token budget based on problem complexity, and Liu et al. [109] make agents budget-aware using a budget tracker to optimize resource usage. While these methods offer efficient LLM reasoning, they are limited to text data. Despite these advancements, IoT systems face challenges due to diverse environments and data types, which can lead to large, computationally expensive prompts that reduce responsiveness. Thus, optimizing prompts to balance the accuracy of reasoning with computational efficiency specifically designed for IoT

environments is essential for seamless LLM integration in resource-constrained environments.

Agentic LLMs. Agentic LLMs extend conventional prompting by enabling planning, tool use, memory, and autonomous decision-making [110]. In IoT systems, Agentic LLMs act as orchestrators over heterogeneous devices and services, discovering capabilities, composing APIs, and executing multistep tasks under resource constraints. Multi-Agent LLM systems distribute roles across specialized agents with domain-specific tools and policies to improve reliability and scalability. Agents can adapt dynamically to various scenarios, analyze multimodal inputs, generate strategies and reason under uncertainty, and engage in self-improvement [111], [112]. Agentic AI and multi-agent LLMs transform LLMs into autonomous orchestrators capable of planning, tool use, and coordinated control. Agentic AI has been used in various applications such as automated prescriptive maintenance in smart manufacturing [112], distributed IoT automation system (LLMind 2.0) [113], spatiotemporal urban data analysis (CityGPT) [114], and industrial IoT system enhancement (LUMEN) [115]. Moreover, using SLMs on edge devices further improves Agentic LLMs by preserving privacy and reducing latency through local deployment [111].

Fine-tuning. Fine-tuning LLMs on task-specific datasets can significantly enhance response time and accuracy. Fine-tuning is also the most common method for adapting LLMs to enterprise and IoT settings [116]. One effective approach is integrating explicit instructions as training examples, which is particularly beneficial in domains with clear procedures, such as healthcare [100], [103]. The self-instruct technique [117] automates instruction tuning using minimal human-annotated data. Parameter-efficient fine-tuning (PEFT) methods, like Adapter Tuning [118] and Low-Rank Adaptation (LoRA) [119], adjust a small subset of parameters, reducing time and computational resources compared to full fine-tuning [120]. While Adapter Tuning may slightly increase latency, LoRA does not. PEFT methods have been applied to various IoT tasks, including synthesizing RF data [121], developing mobile firmware [122], and predicting traffic patterns [123]. Other approaches, such as LoFit [124], aim to enhance interpretability and truthfulness. Test-time scaling is a recent technique that combines reinforcement learning with supervised fine-tuning to improve domain-specific reasoning in pre-trained LLMs [125]. This method uses a small set of carefully selected samples and limits the number of test tokens generated, allowing LLMs to verify and correct their reasoning steps [125]. It can also leverage different architectures to optimize resource use, such as employing depth-recurrent blocks in transformers for efficient computation without requiring special training data, thereby improving reasoning capabilities [126].

Retrieval-Augmented Generation (RAG). Retrieval-Augmented Generation (RAG) is a promising approach

to customize LLMs for specific domain. RAG integrates external knowledge retrieval with generative capabilities, allowing LLMs to combine relevant, up-to-date, and context-specific information without extensive retraining or dedicated fine-tuning, resulting in more accurate responses [127], [128]. RAG can also reduce the computational burden on IoT devices by offloading knowledge retrieval to more capable edge or cloud resources, making RAG a key enabler for scalable and intelligent IoT applications. This approach has been applied in various IoT domains, including real-time environmental monitoring [127], secure data management [129], chat-based security assistants [130], and automated test-case generation [130]. Various optimized RAG techniques have been developed to enhance energy and memory efficiency for mobile devices [128], [131].

Model Evaluation. Evaluating LLMs presents unique challenges, particularly for IoT-specific tasks. Traditionally, LLM evaluation focused on text-based metrics, but recent efforts have shifted towards benchmarks tailored for IoT applications. For instance, eAIBench [122] covers various mobile AI tasks across multiple data modalities, including images, voice, text, and sensor data, specifically designed to assess LLM-based mobile firmware. Another benchmark, SensorQA [132], features human-generated question-answering datasets centered on sensor data, providing rich evaluation context. Beyond output accuracy, there is increasing emphasis on assessing LLM reasoning capabilities, including understanding causal relationships and drawing inferences. The Order-of-Thought framework evaluates reasoning skills based on sensor data by incorporating varying levels of guidance into zero-shot, chain-of-thought (CoT), and prompt-of-thought (PoT) strategies [106]. Fine-tuned LLMs can also serve as evaluators for other models through the LLM-as-a-judge approach, where they are trained on datasets with specific instructions, desired responses, and evaluation criteria [133]. Fine-tuned LLMs can generate prompts and assess responses from other models, providing a structured evaluation mechanism. Alternatively, general-purpose LLMs like GPT-4 can be used for evaluation due to their proficiency in following instructions. By offering a prompt template that includes both the instruction and expected response, these models can effectively evaluate other LLMs' performance in the LLM-as-a-judge approach. This helps ensure LLM robustness and alignment with the requirements of IoT applications. TinyLLM [134] provides evaluation and optimization for agentic small language models (SLMs) used as agents on edge devices.

Truthfulness. A significant challenge in integrating LLMs into IoT systems is hallucination, where LLMs generate false information. This is particularly dangerous in safety-critical domains like healthcare and autonomous systems, as it compromises reliability [42]. Ensuring the accuracy of LLM outputs is crucial for effective decision-making in IoT environments. Truthfulness can be enhanced, for example, by incorporating a bias vector into specific attention

layers [124] or by using contrast decoding to eliminate false information [135]. Additionally, lie and hallucination detection methods are vital for maintaining the integrity of LLM-generated data [136], [137]. Identifying hidden states that encode truth representations can further enhance truthfulness [138], and this technique can also be extended for hallucination detection. For instance, Zhang et al [137] employed prompting techniques to control LLM internal states based on text truthfulness. Another solution is collaborative LLMs, where multiple models cross-validate outputs by referencing each other, refining inferences through peer review and self-tuning mechanisms [139], [140].

Trustworthiness and Ethics. Trustworthiness is a critical concern, especially given the black-box nature of LLMs, which complicates understanding their inference processes [141]. It is an integral part of ethical use, along with other essential qualities for responsible technology deployment, including interpretability, reliability, explainability, safety, fairness, transparency, and resilience [44], [142], as outlined in regulatory frameworks like the EU AI Act and the US AI Act [143]. In IoT systems, fostering trustworthy interactions is paramount, as these devices often cater to individual preferences and require fair and transparent inference capabilities. To enhance trust, integrating clear reasoning instructions into LLM prompts can elucidate the decision-making process, making the underlying data and algorithms more transparent. Additionally, LLM decisions should be summarized and stored for a defined period or presented to end-users in a digestible format [142]. Solutions like audio explanations can further improve accessibility and user experience while promoting the responsible integration of LLMs in IoT systems.

Privacy and Security. LLM integration to IoT ecosystems presents significant privacy and security challenges, alongside unique opportunities for enhanced protection [144]. IoT devices routinely collect and transmit sensitive data, such as health information and location details, which can be further exposed when processed by LLMs, especially if this data is sent to the cloud or third-party services [3], [144]. This data flow increases the risk of breaches, unauthorized access, and inadvertent leaks, particularly given the vast datasets required to train LLMs, which may inadvertently contain or replicate private user information [145], [146]. Furthermore, LLMs introduce new attack surfaces. For example, adversaries can exploit vulnerabilities through data poisoning, adversarial prompts, or model inversion attacks to manipulate device behavior or extract confidential data [147]–[149]. The resource constraints of many IoT devices complicate the deployment of robust protections, as advanced security mechanisms like encryption and anomaly detection may be challenging to implement at scale or in real time [39]. Additionally, employing federated learning or distributed LLM inference introduces further risks, including model poisoning and gradient manipulation [150], [151]. While LLMs can enhance security by enabling intelligent threat detection and

adaptive policy enforcement, their outputs are susceptible to biases and hallucinations, which may undermine trust and reliability [152], [153]. As LLM-driven IoT applications continue to proliferate, it is essential to develop context-aware, resource-efficient, and resilient privacy and security strategies that effectively address both the capabilities and limitations of these technologies.

Concluding Remarks. The challenges of integrating LLMs with IoT are summarized in Table 2. Despite advancements in techniques like prompt engineering, fine-tuning, and knowledge distillation, their effectiveness in real-world IoT applications is context-dependent. For instance, while few-shot prompting can enhance performance, it may introduce latency, making it unsuitable for real-time use. Knowledge distillation can reduce resource demands but may compromise accuracy and adaptability. In safety-critical IoT systems, transparent and reliable decision-making is essential, highlighting the need for ethical considerations and privacy protection. Ultimately, successful deployment of LLM techniques requires balancing resource efficiency, reliability, accuracy, and trustworthiness according to the specific needs of each application. Understanding these challenges and exploring solutions is crucial for effective integration.

V. LLMS AND INTERNET OF THINGS

This survey provides a comprehensive review of the integration of LLMs with IoT across all functional layers of the IoT ecosystem, serving as a reference for researchers and practitioners. We now move to the main part of our survey, structuring the material following different IoT layers. The topics we discuss along with examples of papers addressing them are summarized in Table 3.

A. IoT Devices

LLMs can increasingly operate on IoT devices even if most of them have not been built for this purpose. We next examine ways to improve the integration of LLMs with IoT devices, covering on-device models, distributed architectures, and hardware accelerators.

On-Device LLMs. On-device LLMs are highly optimized versions of large language models specifically designed for embedded devices, such as mobile firmware [122] and LLM-based operating systems for automated vehicles [9]. These models leverage model compression techniques introduced in the previous section, such as pruning and quantization, to fit within the resource constraints of embedded environments [14]. Examples of on-device LLMs that utilize quantization methods include TinyAgent [90] and LLMEdge [13], and more recently Microsoft Phi-3 and Phi-4-mini [160], Gemma3-4B-IT [161], and Qwen 2.5 [162]. Recent advancements, such as Huff-LLM, introduce innovative lossless compression of LLM weights based on Huffman coding, along with a customized decoder, to significantly reduce model size and energy consumption [154]. Also some works harnessing knowledge distillation have been proposed [155],

TABLE 2: Summary of LLM Challenges, their relevance to IoT, and their applications within IoT

Topic	Relevance to IoT	Applications within IoT
Resource Constraints and Sustainability	LLMs require significant computational resources, which many IoT devices lack. High energy consumption poses sustainability challenges.	Optimizing resource utilization is essential for operational efficiency in battery-powered IoT devices. Techniques like model compression and hardware acceleration are crucial.
Model Optimization	Reduces the size and resource demands of LLMs, making them feasible for resource-constrained IoT devices.	Techniques like pruning, quantization, knowledge distillation, and subnetwork modeling applied in edge and on-device deployments, improving performance in IoT applications.
Prompt Engineering	Enhances LLM capabilities tailored to specific data types and contexts.	In healthcare and smart homes to improve LLMs interaction, incorporating contextual information and optimizing for computational efficiency.
Fine-tuning	Adapts LLMs to specific tasks to improve response time and accuracy.	IoT tasks like synthesizing RF data, developing mobile firmware, and predicting traffic patterns, enhancing domain-specific LLMs use cases.
Retrieval-Augmented Generation (RAG)	Integrates knowledge retrieval with LLM capabilities, reducing computational load.	Used in real-time environmental monitoring, secure data management, and automated test-case generation, enhancing accuracy and efficiency.
Model Evaluation	Addresses unique challenges in assessing LLM performance for IoT-specific tasks.	Utilizes benchmarks tailored for IoT applications to evaluate reasoning capabilities and ensure robustness in IoT settings.
Truthfulness	Ensures LLM outputs accuracy, which is critical for decision-making in safety-critical systems.	Techniques for enhancing truthfulness and detecting hallucinations, ensuring data integrity in healthcare and autonomous systems.
Trustworthiness and Ethics	Addresses the ethical implications of using LLMs in IoT, ensuring fair and transparent interactions.	Integrating clear reasoning in LLM prompts and summarizing decisions enhances user trust and promotes responsible technology deployment.
Privacy and Security	Highlights the privacy and security challenges of integrating LLMs with IoT, including data breaches and adversarial attacks.	Advocates for context-aware security measures, such as encryption and anomaly detection, to protect sensitive data and enhance resilience against cyber threats.

TABLE 3: LLMs in IoT layers

IoT Layer	Topics	References
LLMs and IoT Devices	On-Device LLMs	[9], [13], [14], [90], [122], [154]–[162]
	Distributed LLMs	[86], [163]–[165]
	Dynamic Subnetwork Selection	[166]–[171]
	Hardware Accelerators	[172]–[178]
LLMs and IoT Software Engineering	Automating Software Development Processes	[40], [179]–[185]
	Task Planning and Automation	[8], [9], [35], [36], [183], [186]–[190]
LLMs and Sensing	Enhancing Data Collection	[5], [99], [191]–[197]
	Synthetic Data Generation	[121], [198]–[206]
	Context-Aware Sensing	[6], [7], [207]–[209]
	Sensing Application Development	[210], [211]
	Remote Sensing	[212]–[214]
LLMs and IoT Networking	IoT Network Design and Configuration	[8], [39], [51], [80], [215]
	Network Management and Orchestration	[215]–[222]
	Network Analysis	[223]–[227]
	Vulnerability Analysis and Support	[216], [228]–[233]
	Network Topology Management	[216], [234]–[237]
LLMs and IoT Data Processing	Code Generation for Data Processing	[24], [238]–[240]
	Non-Textual Data	[99]–[101], [105], [240]–[257]
	Enhancing Reasoning	[6], [258]–[263]
	Missing Data and Preprocessing	[259], [263]–[267]
	Feature Extraction	[35], [268]–[275]
	Inferences	[99], [106], [192], [260], [276]
LLMs and IoT Privacy and Security	Privacy Management	[3], [277]–[285]
	Computational Privacy	[3], [279], [286]–[291]
	Computational Security	[3], [236], [291]–[301]
LLMs and IoT Human Interaction	Graphical User Interface Agents	[302], [303]
	Enhancing Context Awareness and Adaptation	[26], [243], [304], [305]
	Personalization and User-Centric Interactions	[8], [306]–[310]
	Personal Assistants	[19], [305], [311]

[156]. For instance, MiniLM [157] employs deep self-attention layers to accurately emulate the behavior of larger LLMs with a reduced parameter count. Ray [158] identifies a specific class of sub-billion parameter models, such as SmolLM (135M/360M) and Qwen-0.5B, which are capable

of running on constrained edge hardware. Additionally, Chen et al. emphasize Parameter-Efficient Fine-Tuning (PEFT) methods like Low-Rank Adaptation (LoRA), which freeze most pre-trained parameters and update only low-rank matrices, thereby reducing the memory required for on-device

adaptation by over 70% compared to full fine-tuning [159]. Future directions include integration of on-device LLMs with federation and other distributed learning paradigms to enable collaborative model training. Hardware advancements, such as neuromorphic architectures, may also be able to reduce resource usage and enable more sophisticated LLM capabilities on constrained devices.

Distributed LLMs. Resource requirements can be minimized by collaboratively distributing LLMs across edge and terminal devices [163]. This strategy facilitates the cooperative optimization of performance trade-offs while accounting for available resources. In [164], a collaborative fine-tuning framework is proposed that utilizes a backbone LLM and adapter layers on edge nodes to adapt models according to their specific resource constraints. CoLLM [165] presents a cooperative deployment strategy that integrates tensor parallelization, adaptive load balancing, and an algorithm to determine the optimal number of collaborators for minimizing latency. Similarly, Edge-LLM is designed to optimize both computational and memory requirements for edge devices [86]. To achieve this, Edge-LLM employs a layer-wise compression method along with an adaptive layer tuning method that uses voting to select the best output. Furthermore, Edge-LLM includes a hardware scheduling module to streamline irregular computation patterns arising from diverse quantization, compression, sparsity, and updates to LLM segments. Despite these advancements, resource constraints remain a significant challenge for deploying distributed LLMs in IoT applications. Future directions involve development of algorithms and optimization frameworks for managing resource usage, and tighter integration with fast speed communication technologies to improve communication efficiency between LLMs.

Dynamic Subnetwork Selection. Dynamic subnetwork selection is a specialized optimization technique that leverages the structure of LLMs to activate only specific subsets of the model, known as *subnetworks*. These subnetworks can be distributed across different devices to optimize resource utilization [71]. However, distributing subnetworks can introduce communication overhead and increase the risk of model failures. An alternative approach is the *Mixture of Experts (MoE)* framework, where each subnetwork is treated as an expert. During inference, only a subset of these experts is activated, enabling improved specialization and resource efficiency. MoE is particularly well-suited for agent-based frameworks, where each expert can function as a specialized LLM agent. When the expert are modeled as agents, this framework is also referred to as the mixture of agents framework [171]. The computational complexity of MoE models depends on effectively managing interactions among experts and selecting the appropriate experts for a given input, which can be addressed through parallelization and resource-aware scheduling for expert activation [166], [167], while task allocation techniques can optimize expert selection based on resource constraints and query character-

istics [168]. In contrast, the *Lottery Ticket Hypothesis (LTH)* posits that within a large, over-parameterized model, smaller, well-initialized subnetworks—referred to as “winning tickets”—can be trained to achieve performance comparable to the original model [169], [312]. Identifying these winning tickets allows for the creation of lightweight subnetworks that maintain high accuracy while using fewer resources. LTH can also enhance distributed learning, having been successfully applied to optimize federated learning [170]. However, this approach can introduce security vulnerabilities; for example, LTH-based federated learning systems have been shown to be susceptible to backdoor attacks [313]. Future directions in this field include reducing communication overhead for subnetworks, developing improved mechanisms for identifying optimal subnetworks, and investigating the security and privacy threats associated with these techniques.

Hardware Accelerators. Hardware accelerators are specialized components that enhance the efficiency of neural network operations by enabling simultaneous processing of multiple tasks, which reduces latency and power consumption [172]. The most common type of accelerator is the Field-Programmable Gate Array (FPGA), which can be reprogrammed to create LLM-specific accelerators that optimize data flow and computation for particular models. For example, custom FPGAs have been developed for LLAMA models [173]. FPGAs are especially beneficial for edge devices where multiple models may need to run concurrently. However, the customization process can be labor-intensive and requires specialized design solutions [174], [175]. Alternatively, Application-Specific Integrated Circuits (ASICs), also known as Neural Processing Units (NPUs), are custom chips designed to accelerate AI workloads without being tied to specific models, making them well-suited for IoT integration. They are particularly effective at reducing pre-fill delays, a major source of inference latency [176]. Recent advancements in microcontroller design also aim to enhance LLM operations on existing CPUs through optimized instruction sets and libraries. These usually focus on specific processor families, limiting their broad applicability [177]. While these techniques significantly improve LLM inference speed and efficiency, the diverse architectures of IoT devices and the evolving nature of LLMs necessitate further research. There is also a need to integrate secure processors, privacy mechanisms, and other critical IoT components. For example, integrating FPGAs with trusted execution environments for secure processing is highly challenging [178]. There is also a growing demand for more general accelerators that are capable of delivering consistent performance across various hardware platforms and LLMs.

B. IoT Software Engineering

The integration of Large Language Models (LLMs) into IoT software development is revolutionizing the creation and maintenance of applications for an expanding array of IoT devices. This section explores how LLMs can automate

and enhance various stages of the software development life-cycle, including automating the software development process, and supporting task planning and automation. Note that LLMs naturally have many other uses within software engineering processes. However, these have already been explored extensively by other surveys, such as [314], and in this section we focus specifically on IoT specific software engineering issues.

Automating Software Development Processes. LLMs can significantly enhance the software development process for IoT devices by generating comprehensive textual requirements and use cases based on informal information [40]. They can expedite code generation and fault detection, e.g., by generating dynamic trigger-action programs (TAPs) for home automation and other domains from natural language requests [179]. Other models, such as IoTcoder [180] and LLM4TAP [181], also utilize LLMs for TAP generation. In addition to code generation, LLMs can facilitate automatic validation by utilizing device-specific parameters and feedback from validation to refine the generated code [182]. LLMs can also create customized AI models or agents based on user goals and device constraints [183], [184]. Furthermore, they can be employed for defect localization and bug detection. For example, ChatDL [185] analyzes code updates and defect reports from Industrial IoT system. By enhancing code summarization and ranking changes based on defect relevance, it can localize software errors more accurately. Future research should focus on integrating automated testing approaches to ensure robust system integration and investigating safety protocols that enhance security in IoT software development. Other directions include exploring LLMs in real-time collaborative coding environments and evaluating their potential to improve development efficiency in cross-platform development.

Task Planning and Automation. LLMs excel in activity planning and workflow design, making them valuable for automating large-scale tasks across diverse IoT ecosystems [8]. Their ability to generate and execute code enables automation across devices with varying deployment characteristics, such as differing locations and environments [35], [36], [183]. Research has demonstrated IoT automation through workflow-based approaches, where LLMs act as workflow engines to ensure seamless task execution. This capability is particularly relevant in smart homes and buildings, where LLMs can orchestrate workflows that integrate diverse subsystems that lack shared spatial characteristics. LLM-powered planning can also benefit city-scale autonomous vehicle deployments or guide and enhance task coordination for robots and micro-mobility vehicles [9], [186], [187]. IoT-Together is an example of automated task planning systems [188]. It employs interactive goal management in natural language to identify user goals and generate necessary service code based on current device restrictions. Another example is DELTA, which introduces a task-decomposition method using scene graphs to develop feasi-

ble plans for robots [189]. Enhancing LLMs' planning capabilities through task decomposition and improved reasoning methods, such as test-time inference, can further advance task planning [190]. Despite these advantages, challenges such as latency, processing speed, and output reliability can limit the effectiveness of LLM-based automation in real-time applications [9]. Future research should focus on developing more efficient algorithms to reduce LLM latency without sacrificing output quality. Hybrid models that combine LLMs with edge computing resources may also improve processing speed and reduce overall energy consumption. Moreover, exploring ways to standardize LLM integration across various IoT devices is essential for facilitating interoperability and scalability.

C. IoT Sensing

LLMs have powerful reasoning capabilities that can enhance the collection, processing, and utilization of sensor data. This section examines how LLMs can be used to improve and streamline sensing related tasks in diverse IoT applications.

Enhancing Data Collection. LLMs can enhance the data collection process by developing effective data collection plans that can be transformed into actionable instructions for IoT devices. For example, prompts can guide the formulation of questions regarding data collection and sampling rates tailored to specific contexts, ensuring that IoT devices operate efficiently [99]. Moreover, fine-tuning LLMs can improve sensing tasks by enabling the integration of sensor data with environmental information or domain-specific knowledge. For instance, LightLLM [191] enhances IoT-based light sensing tasks by incorporating task-related knowledge, which is encoded through prompts to facilitate effective data fusion. Additionally, some works, such as CASIT [192], utilize LLM-based agents to preprocess IoT sensor data according to user commands, thereby reducing the volume of data transmitted. Other research focuses on leveraging LLMs for compressing sensor data in specific domains, such as smart transportation [193], employing adaptive sampling and rescaling techniques alongside Chain of Thought prompting. An et al. introduced the IoT-LLM framework, which utilizes Chain-of-thought (CoT) prompting to interpret complex sensor signals—such as ECG and WiFi Channel State Information (CSI) by converting them into descriptive text prompts before processing. This method grounds the LLM in physical reality, significantly improving reasoning accuracy for sensing tasks compared to baseline models [196]. Additionally, multimodal LLMs (MLLMs) can be employed to dynamically control sensor power states, such as disabling power-consuming LiDAR when RGB-D cameras provide sufficient data, thereby optimizing the energy efficiency of the sensing layer [197].

LLMs also offer a potential way to manage the quality of crowdsourcing data by establishing reputation metrics that quantify the reliability of participating IoT devices and ensuring the data they provide is trustworthy and ac-

curate [194]. Additionally, LLMs can be used to create data-driven soft sensors (or proxies) for variables that are difficult to measure directly but can be inferred from other auxiliary sensors, particularly in industrial settings [195]. LLMs can also be used to provide knowledge about the physical world based on their embedded understanding of the environment and IoT sensor data. For example, LLMs can offer information about a user's location by providing sensor data regarding movement [5]. Future directions include optimizing LLMs for diverse sensor types to improve data accuracy and enhance context-awareness. LLM-based data compression also warrants exploration, particularly the preservation of preserve data integrity despite compression. Other potential directions include the generation adaptive prompts based on real-time sensor feedback, and multi-sensor fusion techniques to provide deeper insights into IoT data.

Synthetic Data Generation. LLMs have potential to significantly improve synthetic data generation, alleviating issues with data scarcity and high collection costs [198]. This is particularly important in situations where collecting real-world data is impractical or expensive. Additionally, synthetic data can mitigate privacy concerns by generating information without exposing sensitive personal details. For instance, LLMs have been used to create textual descriptions of motion data [199]. Their ability to articulate human experiences can also aid in generating sensor data for human activity recognition, especially where existing datasets are limited due to variability in real-world scenarios [201]. By describing plausible ways to perform activities, LLMs enhance datasets for training machine learning models, as demonstrated in generating synthetic sensor data for human activity recognition [200] and ground pressure data for various activities [202].

LLMs also can support other domains, such as RF data generation [121]. By combining domain-specific knowledge with reinforcement learning, LLMs can effectively generate general sensor data [198]. In [203], a framework is proposed that uses LLMs to infer desired data features and select domain-specific generative models to enhance the output. GenG [204] adopts a similar strategy, fine-tuning an LLM for sensor data tasks and employing a transformer-based model for data generation. Additionally, LLMs can automate dataset creation for training small language models on devices with limited resources, such as smart homes, enabling them to plan actions based on user goals [205]. Furthermore, LLMs can optimize sample selection from existing benchmarks using adaptive sampling techniques to reduce training costs while maintaining accuracy [206]. Future directions involve refining synthetic data generation techniques to improve dataset realism and diversity, essential for robust machine learning. Integrating LLMs with generative adversarial networks (GANs) could enhance data quality, while evaluating the scalability of these methods across various IoT applications is crucial for assessing their practical impact.

Context-Aware Sensing. LLMs can enhance sensor data by providing additional context. For instance, they have been used to annotate raw IMU sensor data without requiring additional modalities [207], addressing scalability, cost, and privacy concerns associated with human annotators. LLMs also enrich ambient sensor data in smart homes for human activity recognition, allowing models to rely on textual data and facilitating deployment in spaces with varying floor plans [208]. ContextGPT [6] improves human activity recognition by leveraging LLMs to provide context based on additional raw sensor data from wearable devices and smart home sensors, which helps verify the consistency of recognized activities. Similarly, SensorLLM [7] adds textual descriptions of variations and trends in sensor data to aid in activity annotation. Additionally, LLMs can enhance the quality of crowdsourced spatiotemporal data by incorporating high-quality textual data extracted from the web [209]. Future directions involve integrating diverse sensor inputs and ensuring that LLMs can generate real-time contextual insights to support contextual reasoning.

Sensing Application Development. LLMs can provide a programming framework for context-sensing applications based on common sensors [210]. This framework allows natural language to serve as an interface for processing sensor data, enhancing transparency and reducing the need for programming expertise. It also facilitates feedback incorporation for query refinement, simplifying the application development process. Additionally, LLMs can process data collected from ambient sensors in edge devices, enabling real-time emergency detection and providing appropriate interactions and reminders for users [211]. This capability is particularly valuable in domains such as healthcare and elderly care, where timely interventions can significantly impact outcomes. Future directions involve extending LLM programming frameworks to accommodate a broader range of sensor types and applications, and exploring ways to integrate user feedback mechanisms to create more intuitive and user-friendly context-sensing applications.

Remote Sensing. LLM-based agents can be utilized for remote sensing tasks by receiving user commands in natural language, planning the necessary functions, executing these functions to perform the required sensing, and finally delivering the results to the user, as demonstrated in remote imaging [212]. Additionally, LLMs can optimize the balance between sensing and communication in networks, such as unmanned aerial vehicles (UAVs), where both functions share the same signals. By analyzing these signals, LLMs can identify the optimal trade-off between data collection and communication efficiency, ensuring that UAVs effectively gather information while maintaining reliable communication [213]. Furthermore, leveraging multimodal large language models can enhance remote image sensing by improving visual perception and effectively integrating text and images [214]. This capability can lead to richer context and more actionable insights from remote sensing data. Future

directions involve the integration of feedback mechanisms for improved operational efficiency.

D. IoT Networking

As the scale of IoT deployments increases, the complexity of the networks connecting these devices also rises. LLMs can enhance various aspects of IoT networks, including optimizing network configurations, improving resource allocation, and enhancing resilience. This section surveys the role of LLMs in enhancing IoT networks.

IoT Network Design and Configuration. LLMs have potential to tailor network requirements and specifications to align with specific tasks and deployment scales [39], [215]. For instance, the design of on-device and offloading architectures can greatly benefit from LLMs by selecting networking strategies and offloading models that optimize task distribution across the network based on device location, resource availability and other factors [51]. This capability can significantly improve resource allocation and enhance communication efficiency. LLMs can also play a critical role in maintaining network-wide policies. An example is cloud execution where policies govern data communications outside the device [80]. Ensuring that sensitive data generated by prompts is properly handled – either removed or obfuscated – before transmission to the cloud is crucial [8]. However, a notable drawback of this approach is that processing time and communication latency may increase. Future directions involve examination of advanced algorithms and frameworks that can dynamically adapt to changing network conditions and device capabilities.

Network Management and Orchestration. IoT networks support distributed architectures that accommodate heterogeneous devices using various communication technologies. This diversity presents challenges for employing LLMs, particularly due to differences in communication protocols, data formats, and functionalities among devices. Addressing these challenges requires a consistent contextual understanding of device operations and user interactions [216]. LLMs can facilitate frameworks that guide network architecture design, optimizing performance and offering insights into configuration trade-offs, especially in complex networks like 5G/6G and satellite systems [217]. A notable application is intent-based networking, where LLMs can autonomously generate network configurations with minimal external intervention [218], [219]. LLMs can also be used to query and analyze data in real-time from IoT devices and network infrastructure, supporting network orchestration, dynamic task scheduling and other critical tasks [215]. This real-time data analysis can enhance model inference and facilitate the implementation of geographically distributed LLM-based IoT network architectures [220]. Additionally, LLMs can prioritize offloading for devices that are more frequently active [221] or have higher workload demands [222]. Future directions include integration of LLMs with intent-based networking frameworks to enhance adaptive decision-making

processes that optimize performance while ensuring networks effectively respond to real-time changes in conditions and demands.

Network Analysis. By collecting network traffic data and translating it into natural language, custom LLMs can be developed to identify and describe different devices within IoT networks [223]. Such capabilities are essential for managing large-scale IoT systems, where accurately identifying and categorizing devices can significantly enhance network efficiency and responsiveness. As an example, Meyuhas et al. [224] employ an LLM to automatically detect and label IoT devices. This involves extracting and enriching textual features from network traffic, followed by applying machine learning methods for string matching and using LLMs with zero-shot classification to determine device functions and label them based on confidence levels. A caveat in LLM integration is that their use may also introduce new vulnerabilities that can affect network performance or expose configurations. For instance, while speculative decoding can enhance efficiency and accelerate LLM inference, it may lead to input-dependent variations in network packet timing and size [225]. Similarly, LLMs that analyze network data and parameters can inadvertently lead to data leakage that reveals details of the underlying network architecture [226]. Furthermore, the performance of LLMs can be impacted by factors affecting network reliability, such as interference, signal degradation, and inadequate bandwidth, potentially leading to incomplete or inaccurate data being transmitted to the LLM [227]. Future directions involves developing robust mechanisms to mitigate these vulnerabilities while ensuring the effective integration of LLMs in IoT networks.

Vulnerability Analysis and Support. Beyond device identification, LLMs can be leveraged for vulnerability analysis by diagnosing issues, describing them, and providing guidance on corrective actions to improve network resilience and performance [228], [229]. This approach has been successfully applied in optical networks, where AI agents positioned in the control layer manage the physical layer and facilitate interactions at the application layer [230]. LLMs can also optimize communications by suggesting adjustments that enhance data transmission efficiency and increase intelligence within the network [216]. In wireless networks, in-context learning simplifies model training and hyperparameter fine-tuning [231]. LLMs have been employed for large-scale contextual analysis of network data in real-time, enabling dynamic network topology optimization [232]. In the context of 6G mobile communications, LLMs can enhance network performance and security through service personalization and comprehensive operational assessments [233]. Future directions should focus on enhancing the interpretability of LLM-based vulnerability assessments and exploring deployments that bring LLM capabilities closer to end-user IoT devices.

Network Topology Management. Many IoT deployments are highly dynamic, with devices frequently joining or leaving the network. Maintaining accurate information about the network in these environments can be challenging, even for LLMs [234]. However, by continuously examining network usage patterns, LLMs can reduce response times to failures in network architecture, providing timely insights and practical recommendations. For example, analyzing metrics such as packet loss, retransmissions, and communication channel usage using text-based methods helps to create a clearer picture of normal network behavior, customized for specific devices and their usage contexts [235], [236]. This understanding enables targeted corrective actions, such as adjusting transmission protocols, dynamically reallocating bandwidth based on real-time usage, or optimizing routing configurations to enhance overall network performance [237]. LLMs can also bolster network resilience by aiding in failure detection. For instance, they can establish a comprehensive baseline of expected communication patterns, which can be used to identify anomalies in heartbeat responses that may indicate potential node failures [216]. Future directions include enhancing the responsiveness of LLMs in dynamic environments, ensuring they effectively adapt to changes in network topology and device availability.

E. IoT Data Processing

The ongoing expansion of IoT is increasing the volume and complexity of data generated by interconnected devices. LLMs have emerged as powerful tools for processing this data, providing innovative solutions for code generation, data processing, and inference tasks. This section examines the diverse applications of LLMs in IoT data processing.

Code Generation for Data Processing. Early investigations into the utilization of LLMs for code generation have demonstrated that GPT-3 [24], despite not being specifically trained for this purpose, can generate simple programs from Python docstrings [238]. This capability allows LLMs to streamline the generation of code for modifying input datasets, creating meaningful features, preprocessing IoT sensor data, and conducting IoT task reasoning and analytical processes [239]. By utilizing prompts that correspond to specific data-related tasks, such as data visualization, preprocessing, transformation, and analysis, LLMs can produce tailored scripts in various programming languages [240]. Each prompt typically includes a test case that describes the problem, a source file, and clear instructions for code generation. Unlike natural language samples, which are typically evaluated by human evaluators or heuristic algorithms, the generated scripts undergo validation to ensure they run successfully and meet specified requirements. Popular benchmarking frameworks, such as HumanEval [238], measure the functional correctness of these scripts based on their success rate on the first attempt. Future directions include enhancing the robustness of LLM-generated code, integrating feedback mechanisms for iterative improvements, and extending code

generation capabilities for more complex IoT applications, such as real-time data processing and machine learning model deployment.

Non-Textual Data. LLMs were initially designed for processing and generating text, while IoT systems incorporate diverse sensor modalities, including time series, spatial data, voice, and images. Recent advancements have enabled LLMs to better handle various IoT sensor data types, such as spatio-temporal data [241], [242], raw IMU data [99], [101], [243], sensor data from smart homes [105], and time-series data [244]–[249]. Since LLMs can infer from sparse and incomplete data, several methods have been proposed to effectively represent sensor data. One approach is to include raw sensor data within the prompt [250], [251], while another involves transforming spatio-temporal and time-series data into textual representations, such as strings of numerical digits or sequences of tokens, framing time-series forecasting as a next-token prediction task [252], [253]. Additionally, using statistical summaries instead of full raw data can reduce input size and improve response efficiency by minimizing the need for LLMs to process large datasets [100]. More recently, TimeSeriesScientist [257], an LLM-driven agentic framework, leverages four agents for augmentation and preprocessing, self-planning, forecasting and validation, and reporting to provide an interpretable and transparent time-series forecasting process. With the evolution of Large Multimodal Models (LMMs) like GPT-4V(ision) [254], LLMs have also been extended to analyze images through multimodal plugins such as Bing Image Search [255]. This advancement enhances the understanding and inference of the semantics of audio and visual signals. LLMs can also assist in selecting the best analysis tool for data based on user conversations [256]. More recently, Marble [65] utilized multimodal inputs, such as images, videos, and text to generate 3D worlds and provide spatial intelligence functionalities. However, accommodating errors encountered in operations on mixed data types, such as performing arithmetic operations between numbers and strings, remains a challenge [240]. Future directions include improving sensor data representation to enhance fusion, modeling, and reasoning, as well as developing more robust methods for representing and processing mixed data types.

Enhancing Reasoning. To improve the reasoning capabilities of LLMs, additional context regarding time-series inputs and task instructions can be provided in natural language [258]. By employing advanced natural language processing techniques to capture and reason about complex spatio-temporal patterns, LLMs can be adapted through supervised pre-training on non-textual data, followed by task-specific fine-tuning to enhance predictive accuracy in time-series forecasting [258], [259]. More advanced techniques involve integrating domain knowledge into neuro-symbolic models, which can significantly improve activity recognition accuracy [6]. In addition to these strategies, LLMs can construct complex multimodal representations

by fusing data streams from various sensors [260], and incorporating encoders designed for time-series or spatio-temporal data has also shown promise; for instance, Li et al. [261] utilized neural networks to effectively encode temporal dependencies in spatio-temporal data. Moreover, training or fine-tuning LLMs specifically for time-series data proves to be an effective strategy, as demonstrated by Bian et al. [262], who first trained an LLM on patch-based predictions to learn general time-series features across domains before fine-tuning it for multi-patch forecasting. In addition to text-based representations, LLMs can process structured data, such as graphs or tables, through specialized embedding techniques that capture relational and hierarchical information [263]. Future directions include exploring more sophisticated multimodal models that integrate and support diverse data types.

Missing Data and Preprocessing. Missing data is a common and challenging issue in IoT data analysis tasks, and LLMs are increasingly employed to evaluate and impute these missing values [263], [264]. The approach leverages the pre-trained knowledge of LLMs to understand complex patterns and relationships within the data, enabling the generation of accurate imputations for missing values. For instance, LLMs integrated with a graph attention module can effectively capture spatial and temporal dependencies, resulting in more precise imputations for spatiotemporal IoT sensor data [263]. This capability not only improves data quality but also enhances the reliability of subsequent analyses and predictions. Additionally, LLMs can assist in pre-processing time-series data by employing techniques such as padding and tokenization to manage variations in sequence length and create meaningful sequences [265]. To facilitate the integration of information from multiple time scales and improve prediction accuracy, specialized frameworks like LLM4TS have been developed [259]. Furthermore, task-aware reduction methods that focus on reducing input data for LLMs based on downstream tasks [266] are used to enhance LLM performance by prioritizing relevant data. The Flash-Fusion [267] system uses a similar approach to retrieve query-relevant sensor data by using a planner with a semantic understanding of collected data and the user query. Future directions involve extending these frameworks toward advanced imputation techniques capable of handling missing values while also detecting and processing outliers and anomalous data points.

Feature Extraction. Effective data analysis relies on features that accurately characterize the target problem and correlate with desired outputs. LLMs can leverage their pattern learning capabilities to identify meaningful correlations, automating feature generation and facilitating explainability of features [268], [274]. Effective feature extraction requires well-structured prompts that guide the model in extracting relevant features. These prompts typically include task descriptions, expected inputs or outputs, and relevant examples [272]. They can be developed from initial candi-

dates suggested by LLMs [269], guided by chain-of-thought reasoning [270], or constructed based on insights from other prompts [271]. Domain-specific prompts can further refine LLM outputs to generate structured representations, extracting meaningful insights for improved forecasting [275]. For instance, prompts informed by smart home routines can guide LLMs to propose features tailored to user behaviors [35]. Effective feature extraction can also be achieved through few-shot and zero-shot prompting by utilizing the embedded knowledge of pre-trained LLMs [272]. Additionally, LLMs facilitate interactive feature exploration via natural language queries, allowing users to engage intuitively with the data and extract insights without specialized technical knowledge [273]. With continuous learning capabilities, LLMs can adapt to new data and user feedback, enabling the extraction of more relevant and interpretable features. Future directions include exploring the potential for LLMs to collaborate with domain experts in real-time, refining feature extraction processes to ensure that derived features meet practical application needs.

Inferences. LLMs are increasingly utilized for IoT inference, particularly in causal analysis of sensor data to understand how various parameters influence specific outcomes and to improve decision-making [106]. For example, by supplying LLMs with detailed information about data collection methods, categories of human actions, and well-crafted prompts, they can analyze IMU sensor data to recognize common activities such as sleeping, walking, standing, or bicycling [99], [260]. This capability provides a deeper understanding of human behavior and its impact on outcomes. When models for unseen tasks are unavailable, GANs can be combined with LLMs to make inferences without the need to train new models from scratch [276]. Additionally, LLMs enhance data analysis and inference capabilities through natural language commands, allowing users to interact with the system using everyday language for nuanced data exploration. The effectiveness of inference can be further improved through collaborative LLM agents [192], as seamless communication and cooperation among these agents can enhance the understanding of human actions and optimize the operational efficiency of IoT systems. Future directions include investigating the integration of LLMs with real-time feedback loops from IoT devices to continuously refine inference models.

F. IoT Privacy and Security

IoT ecosystems pose significant security and privacy challenges that LLMs can help to address. This section explores the potential LLMs can bring to improve security measures and enhance privacy protection. Security and privacy risks in LLM integration were covered in Section IV.

Privacy Management. The natural language capabilities of LLMs support the integration of privacy policies, user preferences, and regulatory requirements [280]. For example, users can utilize contextualized prompts to customize their

privacy settings and understand the implications of granting permissions. LLMs can be used to offer interfaces that allow users to use natural language to control their interactions with others based on their preferences [277], [278] and specify the sensor data they are willing to share, thereby guiding users in making informed decisions about their data [279]. These tools can also aid in assessing regulatory compliance and extending legal protections, such as those offered by GDPR (General Data Protection Regulation) [280], [285]. LLMs also facilitate the implementation of privacy-enhancing technologies (PETs), which can be integrated as dedicated AI agents for tasks like anonymization and data minimization [281]–[284]. Finally, LLMs can generate tailored educational content to inform users about best practices for data protection, such as changing default passwords or enabling multi-factor authentication [3]. Future directions include the development of user-friendly interfaces and educational tools that empower users to manage their privacy more effectively.

Computational Privacy. LLMs can also facilitate the integration of privacy protection mechanisms into IoT environments. For instance, LLMs can help evaluate k -anonymity to reduce vulnerabilities to re-identification attacks [286], [287], and they can support differential privacy by determining appropriate perturbation mechanisms or the level of noise that should be added [288]–[290]. Similarly, LLMs can automate the de-identification process by analyzing collected data and applying techniques for anonymization. They can generate unique tokens for pseudonymization and suggest methods for generalization or data swapping [287], [291]. Given the vast amounts of data in IoT deployments, identifying effective protection mechanisms is crucial. While LLMs enhance privacy integration, it is important to recognize that training these models often involves sensitive data related to users' behaviors and environments [3], [279]. Thus, the benefits of LLMs must be balanced against the risks of sensitive data leakage. Future directions involve investigating methods to minimize the exposure of sensitive data during LLM training.

Computational Security. Similarly to privacy, LLMs can play a key role in improving security by explaining security principles, providing guidance on enhancement measures, and offering insights into potential threats [3], [292]. For example, guided prompts can facilitate integrity checks on collected data for anomaly detection considering the type of data and historical records [293]. By analyzing data usage patterns, LLMs can help enforce security policies for data access and sharing, and assist in firewall configuration management by advising on application-specific traffic management requirements. LLMs can recommend suitable security mechanisms, including symmetric or asymmetric encryption standards, cryptographic hashing methods, and multi-factor authentication strategies tailored to IoT contexts [291], [294]. LLMs have been used to identify vulnerabilities, such as risks of denial-of-service (DoS) attacks,

by analyzing usage patterns to detect abnormal traffic [295]. Their capabilities can be leveraged for monitoring security threats across IoT layers. For example, LLMs can reduce risks of backdoor attacks by scanning for vulnerabilities and ensuring firmware integrity checks [296], [297]. Another example is traffic analysis to secure data packets during transmission, generating configuration scripts for security protocols like WPA, IPSec, TLS, or VPN [236]. On the other hand, LLMs are vulnerable to modality attacks via prompt-based injection especially when using multimodal sensor data. To reduce the vulnerability, some frameworks such as JailGuard [300] and E^2AT [301], can be used to evaluate the prompts or to enhance the training process. Future directions include exploring the integration of LLMs with real-time threat intelligence solutions, improving the security of federated and distributed learning by identifying potential vulnerabilities in data sharing [315] and evaluating potential vulnerabilities associated with device usage and context, including communication and computation overheads [298], [299].

G. Human Interaction in IoT

LLM capabilities offer numerous opportunities to enhance human interactions in IoT environments. This section explores these interactions, focusing on their contributions to user interface design, context awareness, personalization, and intelligent personal assistants.

Graphical User Interface Agents. LLMs can significantly enhance the design of graphical user interfaces (GUIs) by developing prescriptive models of user interactions and descriptive models of human behaviors [302]. Traditional IoT design often relies on user studies that capture behaviors limited to specific populations, which may not account for diverse or exceptional contexts. For instance, findings from a suburban study may not apply to urban users who leave smart lights on for security. This highlights the need for more inclusive design approaches that consider a wider range of user behaviors. By integrating LLMs with human behavior descriptors, designers can simulate comprehensive human-machine interaction scenarios, including uncommon behaviors [303]. Prompt templates can be created to gather contextual human behavior inputs, such as, "Given this [scenario X], what is [behavior Y] of [user Z]?" in home energy management systems [302]. Future directions involve assessing the effectiveness of these templates across diverse demographics and contexts to refine models and improve real-world applicability.

Enhancing Context Awareness and Adaptation. IoT systems function in dynamic environments where user needs can change rapidly. To adapt effectively and deliver relevant services, LLMs can continuously learn from user interactions to improve system responses. For instance, an LLM can analyze data from a smart thermostat to optimize heating and cooling based on user preferences. Multimodal LLMs like LIMU-BERT [243] and Llama [26] can interpret sensor

data and generate human-like responses to user inquiries about activities tracked by wearables [304] or assess sleep quality using data from smartwatches [305]. A key area for future research is how LLMs can resolve conflicting data from multiple sensors and ensure accuracy in IoT systems. This involves developing algorithms that prioritize sensor data based on context and user preferences, while also addressing the implications for user privacy, usability, and system reliability.

Personalization and User-Centric Interactions. LLMs can enhance user interaction with IoT systems by providing a unified interface that interprets user intents across various devices using natural language processing and voice commands [306]. This simplifies communication, enabling users to control devices seamlessly, such as saying, "Turn off all lights and set the thermostat to 24 degrees," rather than navigating multiple apps. Other studies, such as Ideal-LLM [316], also enhance multilingual adaptation of LLMs for speech-to-text tasks. End-to-End speech LLMs provide more user-friendly interactions by performing automatic speech recognition for various languages and dialects and generating audio responses [317]–[319]. Furthermore, LLMs can personalize interactions by learning individual preferences and behaviors over time [8], potentially helping to make IoT systems more sophisticated and encouraging wider adoption. To improve personalization, LLMs can generate summaries tailored to user preferences, employing reward functions to refine summarization strategies [307]. A multisensory multitask adapter layer can also enhance the integration of diverse data types and tasks, leading to more efficient interactions in IoT systems [308]. However, the autonomy of LLMs in these interactions raises concerns about user control over data and decision-making, potentially diminishing user agency [309]. To mitigate these concerns, initiatives supporting automated, manual, and hybrid interactions modes allow users to select the approach that best suits their needs or to initiate LLM-controlled sequences based on specific requests [310]. Future research should focus on user-centric frameworks to assess the effectiveness of LLM-based interaction modes on user satisfaction and to explore the advantages and disadvantages of LLM-driven personalization.

Personal Assistants. The diversity of device ecosystems and differing user interfaces can create confusion and lead to underutilization as users often struggle with complex commands. LLMs can be integrated with personal data, devices, and services to streamline tasks, such as managing multiple smart home devices, checking device status (e.g., "Is the door locked?"), or identifying issues (e.g., "Which household appliance consumes the most energy?"). Additionally, LLMs can aggregate data from devices and offer insights into their operation. LLMs can be used to schedule tasks based on contextual factors to reduce cognitive load for users and optimize resource usage through dynamic resource allocation and energy efficiency. Personal LLM Agents can extend

Intelligent Personal Assistants by leveraging advanced semantic understanding and reasoning capabilities to enhance user experiences, automate repetitive tasks, and provide personalized recommendations [19]. Similarly, LLM-powered personal health can connect with wearable health monitoring devices to deliver tailored consultation services [305]. This allows bridging the gap between users and healthcare professionals, including virtual and physical doctors [311]. Future directions include evaluating the effectiveness of various LLM architectures in personal assistance and exploring user perceptions and acceptance of these technologies across different IoT application domains, including healthcare, home automation, and personal productivity.

VI. FUTURE DIRECTIONS AND GAP ANALYSIS

Our survey has provided a comprehensive overview of the various ways in which LLM integration is utilized within modern IoT ecosystems. As the research landscape evolves, new challenges emerge, and some existing ones remain unaddressed. In this section, we discuss emerging directions and trends, as well as the open gaps identified by our survey. The trends and their associated challenges across different IoT layers are summarized in Table 4, while the gaps in current research are outlined in Table 5.

A. FUTURE DIRECTIONS AND EMERGING TRENDS

We next outline some high potential directions for future research that warrant further exploration and benefit the broader IoT ecosystem. These are not intended as exhaustive rather as a sample of diverse research areas where further research is needed to improve the integration of LLMs and IoT. For convenience, these topics are also summarized in Table 4.

#1 IoT-Specific LLMs. While current LLMs excel at processing linguistic and visual data, their potential in handling diverse data modalities within IoT ecosystems remains largely untapped. IoT devices generate varied data types, notably spatial and time-series data, which require specialized processing techniques for accurate analysis. For example, spatial data requires an understanding of geometries, whereas time-series data relies on recognizing temporal relationships. Although initial efforts have been made to develop specialized LLM frameworks – as also discussed during our survey – for tasks like geographic reasoning [320], [321] and wearable sensor data analysis [61], the limitations of LLMs in processing non-textual and non-visual data types are still not well understood. Further research is essential to identify these limitations and to create algorithms and models tailored specifically to these data types. There is also a need for IoT-specific foundation models that support various data modalities. This requires curating datasets that accurately reflect the unique characteristics of IoT environments, addressing storage, processing, security, and privacy needs.

#2 Penetrative AI. Penetrative AI extends LLM capabilities to interpret and articulate real-world dynamics through con-

TABLE 4: Challenges of Integrating LLM Technologies into IoT Layers

Challenges	Devices	Software Engineering	Sensing	Networking	Data Processing	Privacy & Security	Human Interaction
#1. IoT-Specific LLMs	Processing architectures for non-textual data	Development of specialized LLM frameworks	Understanding spatial and time-series data	Network variations in data transfer	New algorithms for diverse data types	Concerns with sensor data integration	User-friendly interfaces for customization
#2. Penetrative AI	Advanced HW for continuous sensing	Integration of LLMs with existing software	Improved sensing technologies	Robust networking for real-time data	Data processing for massive data	Real-time data privacy and security concerns	Semantic navigation interfaces
#3. High-Level Reasoning & Cognitive IoT	Adaptation for collaborative reasoning	Support for interoperability	Enhanced internal/external state sensing	Connectivity for collaborative reasoning	Real-time adaptability in data analysis	Security in collaborative environments	User interfaces for complex interactions
#4. LLM-Enabled Semantic Comms	Hardware support for LLMs	Model training and adaptation	Sensing for accurate data interpretation	Optimization for reduced data transmission	Ensuring semantic interpretation	Security in semantic communications	Interfaces for agent-based communication
#5. Quantum-Enhanced LLMs	Quantum-IoT Hardware Integration	Software for quantum-enhanced models	Support to quantum sensors	Networking for quantum data transfer	Complexity in managing large data volumes	Security of quantum-enhanced models	Quantum-Enhanced UIs
#6. Human-Robot Interactions	Hardware for advanced robotics	LLMs for understanding natural language	Sensing for contextual awareness	Networking for real-time communication	Analysis of user interactions	Ethics and safety in robotic operations	Action languages for intuitive interactions
#7. Secure Execution and LLMs	TEEs for secure execution of LLMs	LLMs adaptation for resource constraints	Multi-device sensing implementation	Network partitioning for model execution	Efficient models for constrained environments	Protecting sensitive data in LLMs	Efficient interaction with secure models
#8. Hyperautomation	Hardware for supporting hyperautomation	LLM integration in automation software	Sensor LLM sampling coordination	Networking for seamless automation	Contract validation and compliance	Safeguarding against legal risks	Natural language interfaces for contract creation

tinuous environmental sensing. By creating digital representations akin to digital twins, opportunities arise for innovative semantic applications. For example, semantic localization can identify specific IoT deployments [322], while semantic navigation can utilize IoT maps for enhanced wayfinding [323]. Combining these capabilities with LLMs facilitates intuitive interactions, e.g., through the dynamic generation of semantic navigation maps using natural language. At the same time, achieving this vision requires advances across all IoT layers. While current sensing technologies effectively characterize people and objects, further advances are needed to detect changes in internal and external states. By constructing behavioral profiles and aligning them with the emotional state characteristics of users [324], we can enhance interactions between humans, digital environments, and the physical world, fostering a more seamless and intuitive integration [325]. Achieving high-fidelity representations requires integrating diverse sensors, which may be distributed and generate massive amounts of data. Efficiently processing this data is a considerable challenge, as is coping with network variations that can affect data transfers. Additionally, the use of real-time data raises significant privacy and security concerns.

#3 High-Level Reasoning and Cognitive IoT. IoT ecosystems are characterized by a high degree of heterogeneity in operating environments, devices, users, and software. LLMs are enhancing the integration of these devices by improving interoperability, which previously required significant effort

to adapt solutions from one domain to another [326]. Furthermore, the advancing reasoning capabilities of LLMs enable the generation of high-level contextual descriptions that individual data sources alone cannot provide [241], [260]. This capability allows for a more nuanced understanding of context, reducing the reliance on tailored domain-specific rules or descriptions that were essential prior to the advent of LLMs [327]. While these developments showcase LLMs' potential to enhance comprehension, these works have also highlighted significant variations in the reasoning support offered by different models [260]. These discrepancies call for further research to explore how LLM-based reasoning can be better integrated into IoT applications, particularly in enabling real-time adaptability to changing conditions and user needs. Such adaptability is vital for ensuring that IoT systems can react to dynamic environments. Another promising avenue is the establishment of reasoning capabilities across a network of IoT devices. This interconnectedness could facilitate the emergence of a Cognitive IoT, where devices collaboratively reason and analyze the data they collect. By fostering collaboration among IoT devices, these systems can engage in complex decision-making processes leveraging shared knowledge and insights. Such advancements would significantly enhance the functionality and intelligence of IoT applications, enabling them to operate more autonomously and effectively.

#4 LLM-Enabled Semantic Communications. Future IoT communication systems are expected to prioritize semantic

communication where the focus shifts from raw data transmission to meaning and context. This approach reduces the need to send large data streams, alleviating demands on the network. LLM-enabled communications extend this by leveraging LLMs to support data interpretation and bridge semantics across communication end points [328]. Semantic communication is expected to be particularly useful for agent-based approaches where it can enhance both human-agent and agent-agent communication [329]. While semantic communication can leverage contextual understanding and reasoning capabilities of LLMs, the limitations of the knowledge model also control what the communications can offer [330]. There have been some efforts to enable LLM-enabled semantic communication, such as LLM-Twin [330] which targets semantic communication for digital twins. However, ensuring the accuracy and reliability of semantic interpretation across diverse devices and domains requires robust model training and continual adaptation. Similarly, to leverage LLMs, we need to deploy them on all devices, which depends on advances in model compression and edge computing [330]. Despite these challenges, the potential of semantic communications is vast and it can enable a new generation of more intelligent and context-aware networks.

#5 Quantum-Enhanced LLMs. Quantum computing offers a promising pathway to achieving superior performance and facilitating large-scale learning of complex machine learning models, including LLMs [331]. Although current quantum circuits are not yet sufficiently mature for widespread, mainstream application, there are already opportunities to leverage existing quantum circuits to enhance LLM performance. For example, the weight matrices in the deep layers of LLMs can be replaced with variational quantum circuits, potentially reducing the memory footprint of LLMs without compromising accuracy [332]. Despite the exciting possibilities that quantum-enhanced LLMs present, existing research has primarily focused on improving or accelerating computations for current LLM architectures, rather than exploring the integration of quantum circuits into IoT ecosystems. This gap highlights an opportunity for further investigation into how quantum computing can support the deployment and integration of LLMs within IoT frameworks. The potential applications of combining LLMs and quantum computing in IoT ecosystems are vast, particularly in scenarios involving the processing of large volumes of data or managing complex models characterized by a high number of parameters.

#6 Human-Robot Interactions. Robotics platforms are increasingly becoming integral components of our daily lives, serving as specialized examples of IoT systems. They are already readily utilized in a variety of real-world applications, including delivery robots, manufacturing automation, healthcare assistance, and domestic support. Moreover, the emergence of advanced humanoid robots that are capable of mimicking human gestures, behaviors, and expressions—signals a significant leap in capabilities. LLMs are seen as a potential enabler for next-generation robotics platforms [333].

These models empower robots to process and understand natural language, which enables users to communicate their needs and preferences seamlessly. LLMs integration can also enhance a robot's ability to interpret contextual cues and emotional subtleties. This leads to interactions that can be more empathetic and highly responsive. As robots become adept at grasping nuanced language, they can adjust their behaviors in real-time, providing personalized assistance. As the field continues to evolve, it is crucial to address the challenges associated with LLM integration to ensure that robotic systems are safe, ethical, and beneficial to society. Key considerations include the development of techniques for hallucination detection and critical threat recognition, which are essential for maintaining safe robotic operations. There is also a need to create more descriptive action languages that provide clear behavioral models for robots to follow. Early research indicates that such descriptions can significantly enhance the responsiveness of robotic platforms, enabling complex tasks like the autonomous manipulation of objects [334]. Furthermore, the successful integration of LLMs into robotics necessitates ongoing advancements in system design and LLM execution to guarantee that interactions with LLMs are fast and resource-efficient.

#7 Secure Execution and LLMs. Commercial LLMs typically operate within data centers or other infrastructure that is physically secure. This contrasts sharply with IoT devices that can be deployed in unsecured environments, making them vulnerable to physical access by malicious actors and raising concerns about the security of both the data and the AI models on these devices [335]. To address these challenges, IoT devices increasingly integrate Trusted Execution Environments (TEEs), which provide tamper-resistant execution through secure cryptoprocessors and encryption. However, adapting LLMs to run within TEEs is complex due to their limited memory and processing capabilities, which can hinder the efficient execution of large models [336]. A promising approach involves identifying sensitive components of the model that require heightened protection, allowing these parts to run in the TEE while leveraging external GPU and CPU resources for less sensitive components [337]. Nonetheless, this may lead to potential leakage of sensitive information, such as key-value (KV) cache values, unless adequately protected [338]. Another solution is to partition the model across a network, enabling each device to execute only the portions that fit within their TEEs. Advances in LLM efficiency are crucial for adapting LLMs for TEE execution as they facilitate model size reduction and improve execution efficiency in resource-constrained environments. Continued research is essential to enhance the safety of LLMs and the data they process in IoT-rich environments, ensuring secure deployment.

#8 Hyperautomation. LLM integration into IoT ecosystems can support scaling up automation toward *hyperautomation* where business processes and workflows are heavily automatized [339]. While there currently are technologies for

automating business processes, such as the use of smart contracts that operate on top of Blockchain or another distributed ledger technology [340], these remain somewhat limited due to the complexity in establishing the contracts and verifying them. LLM technology can facilitate and automate contract creation, offering natural language interfaces for specifying the conditions that the contracts should satisfy, and help translate the requirements into executable code. This streamlines the creation process, and decreases the level of expertise that is needed for creating them. The advanced reasoning capabilities of LLMs are also important for large-scale automation as they facilitate validating contracts and support conflict resolution. The validation capabilities can also be extended to ensure continued compliance with legal and regulatory frameworks, safeguarding industries against legal and compliance risks and minimizing the need to recreate contracts as regulatory frameworks change. While the benefits of LLMs to scaling up workflow automation are clear, further research is needed to better integrate LLMs with current automation processes. Among others, this involves creation of natural-language interfaces for supporting contract creation, establishment of adversarial models that can be used to analyze pitfalls of LLM-based contract validation, as well as improved LLM reasoning capabilities that can support contract validation. There is also a need for studying in specific fields, such as smart grid management or logistics, on how LLMs and automation processes can be further enhanced.

TABLE 5: Mapping of Research Gaps to IoT Layers

Challenges (Research Gaps)	Device	Software Engineering	Sensing	Networking	Data Processing	Privacy & Security	Human Interaction
1. LLM Optimization for Diverse Sensors	✓		✓		✓		
2. Real-Time Contextual Insights		✓	✓				
3. Integration of Feedback					✓		✓
4. User Feedback Mechanisms in Applications		✓					✓
5. Privacy and Security Risks Assessment		✓				✓	
6. Integration with Privacy Mechanisms						✓	✓
7. Gaps in Practical Implementations	✓			✓			
8. Cross-Domain Adoption		✓		✓			
9. Real-Time Feedback Loops for Inference Models		✓			✓		
10. Missing Data Handling			✓				
11. LLM Interoperability and Scalability	✓						✓
12. Ethical Considerations in LLM Deployment			✓			✓	✓

B. GAP ANALYSIS

LLM integration into IoT ecosystems presents a transformative opportunity to enhance every aspect of IoT ecosystems. Yet, as we delve deeper into the current landscape, we can identify several challenges that remain unaddressed. Below, we discuss these research gaps in detail, summarizing their implications and the need for further investigation, as also illustrated in Table 5.

#1 LLM Optimization for Diverse Sensors. Current software and hardware optimization techniques primarily focus on enhancing the efficiency of existing LLM techniques or specific model types on IoT devices. However, as discussed in the previous section, there is also a need to develop LLMs that can effectively handle other data types, such as time series sensor data or spatial data. These data types may necessitate models with different architectures or algorithms, and existing acceleration and optimization techniques may not be optimal for these specialized models. There is also a need for establishing benchmarks on other data types, and carrying out evaluations assessing the potential and the limitations of current optimization and acceleration techniques.

#2 Real-Time Contextual Insights. The ability to analyze multi-sensor inputs and generate real-time insights about the operating environment is crucial for enabling LLMs to support context-awareness across different applications. Achieving this capability demands a tight integration of LLMs into applications designed to process and analyze sensor data. Currently, there is limited support for such integration. Software frameworks that can facilitate this process would be highly beneficial for ensuring seamless integration. At the same time, there is also a need to support the adaptation of models to dynamically changing environments. This may involve, e.g., detecting data drift and deploying agents that focus on specific sensors or sensor combinations as the available data evolves.

#3 Integration of Feedback. LLMs increasingly integrate mechanisms that allow them to learn and evolve over time. Current techniques are well-suited for data where the relationship between input and output is relatively straightforward, such as textual or visual data. How to bridge these models for more complex process flows, such as complicated IoT data-analysis pipelines, remains an open challenge. Further research is needed to explore methods for effectively capturing and utilizing feedback to enhance data analysis and decision-making processes.

#4 User Feedback Mechanisms in Applications. IoT applications operate in a wide range of contexts and on diverse devices. These devices may lack traditional input capabilities or the operating contexts may be such that providing feedback is difficult or even dangerous. Fields such as recommender systems and virtual reality have explored implicit feedback methods [341], [342], such as capturing and analyzing user gaze patterns or monitoring interaction durations. Similar mechanisms are needed for LLMs integrated into IoT applications operating in complex real-world settings, such as industrial manufacturing, healthcare, mining, or military applications. Further research is needed to develop user-friendly methods for gathering feedback and to seamlessly integrate these mechanisms into the LLM life-cycle. There is also a need to evaluate different mechanisms for gathering and integrating implicit feedback with LLMs. By doing so, we can refine and evolve model outputs as

well as enhance both the effectiveness and user experience of LLM-powered IoT solutions.

#5 Privacy and Security Risks Assessment. The heterogeneity of data and devices in IoT environments makes ensuring privacy and security challenging. Foundation models or large language models trained on diverse IoT data could support auditing these systems and identifying potential vulnerabilities. They can analyze real-time data streams to detect outliers, such as targeted poisoning attacks or other anomalies, and generate security reports that highlight weaknesses. Additionally, LLMs can help organizations understand user privacy concerns by analyzing interactions, particularly in sensitive areas like healthcare. However, the integration of LLMs also introduces privacy risks, such as data leakage and prompt injection, which can compromise system security [343].

To maximize the benefits while mitigating these risks, strong governance frameworks are essential. This necessitates further research into best practices for dataset curation, regular audits, and the ethical use of these models to ensure robust privacy and security in IoT environments.

#6 Integration with Privacy Mechanisms. In many IoT applications, sensor data may appear innocuous but can inadvertently expose sensitive information when analyzed by LLMs. This is particularly concerning in contexts where seemingly benign data can be combined with other datasets to reveal personal or confidential information. To minimize the risk of compromising sensitive information during analysis, it is essential to support robust interactions between LLMs and established privacy mechanisms. For instance, LLMs should be capable of performing meaningful analyses on differentially private data or k-anonymized datasets, ensuring that individual privacy is preserved while still deriving valuable insights. Furthermore, additional research is needed to prevent the leakage of sensitive information in LLM outputs and to integrate privacy mechanisms directly into the outputs generated by LLMs. Similarly, there is need to integrate with privacy management interfaces that can enable configuring which data types to consider and which to exclude.

#7 Gaps in Practical Implementations. While the diverse frameworks for LLM integration into IoT ecosystems are promising, significant gaps remain in deploying them in practical implementations. IoT ecosystems integrate many legacy systems that are not designed to accommodate advanced AI/ML techniques and even more recent hardware may not be suitable for effective LLM deployments unless specifically designed for this purpose. Different subsystems of a larger IoT deployment can also be heavily fragmented and leverage different transmission protocols, security mechanisms, and other components. Further research is needed on establishing ways to facilitate integration of LLM-based frameworks into diverse IoT ecosystems, including legacy solutions. Similarly, there is need for user-friendly frame-

works that enable compressing and fine-tuning models for specific enterprise scenarios, facilitating the practical adoption of LLMs across different IoT ecosystems.

#8 Cross-Domain Adoption. The emergence of agent-based solutions has the potential to further boost the adoption of specialized models that are tailored for specific application domains. Enabling cross-domain interactions between models and agents, or chaining agents, can potentially enable new applications and further boost the capabilities of LLMs integrated into IoT ecosystems. Harnessing this potential requires a better understanding of the potential and the limitations of domain-specific LLM models, as well as protocols and frameworks that can enable cross-domain interactions.

#9 Real-Time Feedback Loops for Inference Models. Real-time feedback loops that enable continuous adaptation of LLM parameters based on immediate data and user interactions are critical for ensuring accurate and relevant responses in dynamic IoT environments. These feedback loops are particularly vital in safety-critical scenarios, where timely analysis of information is essential and the accuracy of responses can significantly impact outcomes. For instance, IoT tools designed to assist emergency responders must provide instant access to relevant information, while factory workers depend on precise guidance to maintain safety and efficiency. Implementing effective feedback loops requires algorithms capable of rapid data stream analysis to transform this data into actionable inputs for the dynamic adjustment of LLM parameters as conditions evolve. By facilitating continuous learning and immediate adaptation, these feedback loops can greatly enhance the utility and effectiveness of LLMs in IoT applications, ensuring they remain responsive to the dynamically changing needs of users and their environments.

#10 Missing Data Handling. Handling missing data and anomalies in sensor data is a significant challenge in IoT data analysis [344]. The reasoning capabilities of LLMs could assist in this task by offering new ways to infer missing values and manage incomplete, corrupted, or erroneous data. While dedicated imputation techniques may prove more effective than LLMs in many scenarios, their effectiveness typically varies based on the mechanisms leading to data loss. Consequently, even if LLMs are unable to provide effective imputation techniques, they can still offer insights into these mechanisms and help select the most suitable imputation technique for a given situation.

#11 LLM Interoperability and Scalability. As paradoxical as it may seem, while LLMs can enhance the interoperability of data and analyses across different IoT devices, their increasing deployment also necessitates ensuring that LLMs are interoperable with one another. Particularly when a large number of smaller LLMs (or SLMs) operate in tandem, it is essential to ensure that their combined responses are coherent and free from hallucinations. Simultaneously, there is a need to improve scalability by facilitating the integration of LLMs

onto various platforms, such as by customizing the number of parameters to better fit the model on a given device.

#12 Ethical Considerations in LLM Deployment. While ethical guidelines for AI use are gradually emerging [345], many IoT scenarios fall outside their scope, particularly when previously unconsidered inputs are involved. The chaining and linking of AI models, including LLMs, further exacerbate concerns, as even benign individual inferences may enable reasoning about highly intrusive information. For example, while the amount of physical activity and the radius of movement are relatively innocuous attributes on their own, together they can provide cues about an individual's socioeconomic status, leading to unwarranted uses of personal information. Although the privacy of sensor data is actively researched, there is a critical need to explore the ramifications of integrating IoT sensor data with LLMs. To ensure ethical operations, future research should focus on understanding these implications and developing robust guidelines to protect user privacy.

VII. CONCLUSIONS

The integration of LLMs with IoT represents a transformative development in how we manage, interpret, and interact with vast networks of interconnected IoT devices. This survey has systematically explored the roles that LLMs can play across various layers of IoT architecture, addressing challenges and opportunities presented by this convergence. As the number of connected IoT devices continues to grow, the complexity of data management and contextual interactions becomes increasingly pronounced. Traditional processing methods are often inadequate, underscoring the necessity for advanced frameworks that can harness the capabilities of LLMs.

Our comprehensive review of over 300 articles has described the numerous ways in which LLMs can enhance IoT functionalities. From improving data interpretation and enabling natural language interactions to automating tasks and facilitating data-driven decision-making, the potential applications are vast. However, the successful integration of LLMs into IoT systems is not without its challenges. Issues such as resource constraints, latency, and privacy concerns must be navigated carefully to maximize the benefits of LLMs while minimizing associated risks. The findings of our survey highlight that the integration of LLMs in IoT has made significant strides in understanding how these technologies can work together, yet many challenges and opportunities remain for the future. We have identified critical research gaps that warrant further exploration, including the need for specialized LLM architectures tailored to diverse IoT data types, effective integration strategies, and robust privacy and security mechanisms. Additionally, emerging technologies such as quantum computing present exciting avenues for enhancing the capabilities of LLMs in IoT environments.

Looking forward, ongoing research and innovation is essential to unlock the full potential of LLMs in enhancing IoT ecosystems. Our comprehensive survey serves as a foundational resource for researchers and practitioners alike, providing insights that can drive meaningful advancements in the intersection of LLMs and IoT, leading to improvements across all aspects of IoT ecosystems and enabling new applications and other opportunities.

REFERENCES

- [1] J. Ding, M. Nemati, C. Ranaweera, and J. Choi, "IoT connectivity technologies and applications: A survey," *IEEE Access*, vol. 8, pp. 67 646–67 673, 2020.
- [2] Statista. (2024) Number of Internet of Things (IoT) connections worldwide from 2022 to 2023, with forecasts from 2024 to 2033. Accessed on 04 September 2024. [Online]. Available: <https://www.statista.com/statistics/1183457/iot-connected-devices-worldwide/>
- [3] Y. Yao, J. Duan, K. Xu, Y. Cai, Z. Sun, and Y. Zhang, "A survey on Large Language Model (LLM) security and privacy: The good, the bad, and the ugly," *High-Confidence Computing*, p. 100211, 2024.
- [4] S. Minaee, T. Mikolov, N. Nikzad, M. Chenaghlu, R. Socher, X. Amatriain, and J. Gao, "Large Language Models: A survey," *arXiv preprint arXiv:2402.06196*, 2024.
- [5] H. Xu, L. Han, Q. Yang, M. Li, and M. Srivastava, "Penetrative AI: Making LLMs comprehend the physical world," in *Proceedings of the International Workshop on Mobile Computing Systems and Applications*, 2024, pp. 1–7.
- [6] L. Arrotta, C. Bettini, G. Civitarese, and M. Fiori, "ContextGPT: Infusing LLMs knowledge into neuro-symbolic activity recognition models," in *IEEE International Conference on Smart Computing (SMARTCOMP)*. IEEE, 2024, pp. 55–62.
- [7] Z. Li, S. Deldari, L. Chen, H. Xue, and F. D. Salim, "SensorLLM: Aligning Large Language Models with motion sensors for human activity recognition," *arXiv preprint arXiv:2410.10624*, 2024.
- [8] H. Wen, Y. Li, G. Liu, S. Zhao, T. Yu, T. J.-J. Li, S. Jiang, Y. Liu, Y. Zhang, and Y. Liu, "Autodroid: LLM-powered task automation in Android," in *Proceedings of the Annual International Conference on Mobile Computing and Networking*, 2024, pp. 543–557.
- [9] J. Ge, C. Chang, J. Zhang, L. Li, X. Na, Y. Lin, L. Li, and F.-Y. Wang, "LLM-based operating systems for automated vehicles: A new perspective," *IEEE Transactions on Intelligent Vehicles*, 2024.
- [10] "Amazon Integrates Alexa into BMW Cars with Teams from Qualcomm and Auto AI, author=Chain Store Age," *Chain Store Age*, 2025, accessed: 2025-08-11. [Online]. Available: <https://chainstoreage.com/amazon-integrates-alexa-bmw-cars-teams-qualcomm-auto-ai>
- [11] PepsiCo, "Artificial Intelligence at PepsiCo," 2025, accessed: 2025-08-11. [Online]. Available: <https://www.pepsico.com/our-stories/story/artificial-intelligence-at-pepsico>
- [12] Microsoft, "Machine learning inference on iot edge," 2025, accessed: 2025-08-11. [Online]. Available: <https://learn.microsoft.com/en-us/azure/architecture/guide/iot/machine-learning-inference-iot-edge>
- [13] P. P. Ray and M. P. Pradhan, "LLMEdge: A novel framework for localized LLM inferencing at resource constrained edge," in *International Conference on IoT Based Control Networks and Intelligent Systems (ICICNIS)*. IEEE, 2024, pp. 1–8.
- [14] J. Xu, Z. Li, W. Chen, Q. Wang, X. Gao, Q. Cai, and Z. Ling, "On-device language models: A comprehensive review," *arXiv preprint arXiv:2409.00088*, 2024.
- [15] S. Zhang, L. Dong, X. Li, S. Zhang, X. Sun, S. Wang, J. Li, R. Hu, T. Zhang, F. Wu, and G. Wang, "Instruction tuning for Large Language Models: A survey," *arXiv preprint arXiv:2308.10792*, 2023.
- [16] S. Vatsal and H. Dubey, "A survey of prompt engineering methods in Large Language Models for different NLP tasks," *arXiv preprint arXiv:2407.12994*, 2024.
- [17] G. Bai, Z. Chai, C. Ling, S. Wang, J. Lu, N. Zhang, T. Shi, Z. Yu, M. Zhu, Y. Zhang, X. Song, C. Yang, Y. Cheng, and L. Zhao, "Beyond efficiency: A systematic survey of resource-efficient Large Language Models," *arXiv preprint arXiv:2401.00625*, 2024.
- [18] H. Naveed, A. U. Khan, S. Qiu, M. Saqib, S. Anwar, M. Usman, N. Barnes, and A. Mian, "A comprehensive overview of Large Language Models," *arXiv preprint arXiv:2307.06435*, 2023.

- [19] Y. Li, H. Wen, W. Wang, X. Li, Y. Yuan, G. Liu, J. Liu, W. Xu, X. Wang, Y. Sun, R. Kong, Y. Wang, H. Geng, J. Luan, X. Jin, Z. Ye, G. Xiong, F. Zhang, X. Li, M. Xu, Z. Li, P. Li, Y. Liu, Y.-Q. Zhang, and Y. Liu, "Personal LLM agents: Insights and survey about the capability, efficiency and security," *arXiv preprint arXiv:2401.05459*, 2024.
- [20] L. Wang, C. Ma, X. Feng, Z. Zhang, H. Yang, J. Zhang, Z. Chen, J. Tang, X. Chen, Y. Lin, W. X. Zhao, Z. Wei, and J. Wen, "A survey on Large Language Model based autonomous agents," *Frontiers of Computer Science*, vol. 18, no. 6, p. 186345, 2024.
- [21] E. Ferrara, "Large Language Models for wearable sensor-based human activity recognition, health monitoring, and behavioral modeling: A survey of early trends, datasets, and challenges," *arXiv preprint arXiv:2407.07196*, 2024.
- [22] X. Zhang, R. R. Chowdhury, R. K. Gupta, and J. Shang, "Large Language Models for time series: A survey," *arXiv preprint arXiv:2402.01801*, 2024.
- [23] B. D. Lund and T. Wang, "Chatting about ChatGPT: How may AI and GPT impact academia and libraries?" *Library Hi Tech News*, vol. 40, no. 3, pp. 26–29, 2023.
- [24] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. M. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, "Language models are few-shot learners," *Advances in Neural Information Processing Systems*, vol. 33, pp. 1877–1901, 2020.
- [25] R. Thoppilan, D. De Freitas, J. Hall, N. Shazeer, A. Kulshreshtha, H.-T. Cheng, A. Jin, T. Bos, L. Baker, Y. Du, Y. Li, H. Lee, H. S. Zheng, A. Ghafouri, M. Menegali, Y. Huang, M. Krikun, D. Lepikhin, J. Qin, D. Chen, Y. Xu, Z. Chen, A. Roberts, M. Bosma, V. Zhao, Y. Zhou, C.-C. Chang, I. Krivokon, W. Rusch, M. Pickett, P. Sriniivasan, L. Man, K. Meier-Hellstern, M. R. Morris, T. Doshi, R. D. Santos, T. Duke, J. Soraker, B. Zevenbergen, V. Prabhakaran, M. Diaz, B. Hutchinson, K. Olson, A. Molina, E. Hoffman-John, J. Lee, L. Aroyo, R. Rajakumar, A. Butryna, M. Lamm, V. Kuzmina, J. Fenton, A. Cohen, R. Bernstein, R. Kurzweil, B. Aguera-Arcas, M. Croak, E. Chi, and Q. Le, "LaMDA: Language models for dialog applications," *arXiv preprint arXiv:2201.08239*, 2022.
- [26] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, and A. Rodriguez, "Llama: Open and efficient foundation language models," *arXiv preprint arXiv:2302.13971*, 2023.
- [27] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2018.
- [28] L. Sanna, P. Bellan, S. Magnolini, M. Segala, S. G. Haez, M. Consolandi, and M. Dragoni, "Building certified medical chatbots: Overcoming unstructured data limitations with modular rag," in *Proceedings of the First Workshop on Patient-Oriented Language Processing (CL4Health)@ LREC-COLING 2024*, 2024, pp. 124–130.
- [29] S. Mukherjee, P. Gamble, M. Sanz Ausin, N. Kant, K. Aggarwal, N. Manjunath, D. Datta, Z. Liu, J. Ding, S. Busacca, C. Bianco, S. Sharma, R. Lasko, M. Voisard, S. Harneja, D. Filippova, G. Meixiong, K. Cha, A. Youssefi, M. Buvanesh, H. Weingram, S. Bierman-Lytle, H. S. Mangat, K. Parikh, S. Godil, and A. Miller, "Polaris: A safety-focused LLM constellation architecture for healthcare," *arXiv preprint arXiv:2403.13313*, 2024.
- [30] R. Wu, C. Yu, X. Pan, Y. Liu, N. Zhang, Y. Fu, Y. Wang, Z. Zheng, L. Chen, Q. Jiang, X. Xu, and Y. Shi, "MindShift: Leveraging Large Language Models for mental-states-based problematic smartphone use intervention," in *Proceedings of the CHI Conference on Human Factors in Computing Systems*, 2024, pp. 1–24.
- [31] C. Xiao, S. X. Xu, K. Zhang, Y. Wang, and L. Xia, "Evaluating reading comprehension exercises generated by LLMs: A showcase of ChatGPT in education applications," in *Proceedings of the 18th Workshop on Innovative Use of NLP for Building Educational Applications (BEA 2023)*, 2023, pp. 610–625.
- [32] G. M. Fernández-Nieto, V. Echeverria, R. Martinez-Maldonado, and S. B. Shum, "Yarnsense: Automated data storytelling for multimodal learning analytics," in *Data Storytelling and Learning Analytics Workshop 2024*. CEUR-WS, 2024, pp. 124–138.
- [33] D. J. Samuel, Y. Sermet, D. Cwiertny, and I. Demir, "Integrating vision-based AI and Large Language Models for real-time water pollution surveillance," *Water Environment Research*, vol. 96, no. 8, p. e11092, 2024.
- [34] R. Chandra, S. S. Kumar, R. Patra, and S. Agarwal, "Decision support system for forest fire management using ontology with Big Data and LLMs," *arXiv preprint arXiv:2405.11346*, 2024.
- [35] E. King, H. Yu, S. Lee, and C. Julien, "Sasha: Creative goal-oriented reasoning in smart homes with Large Language Models," *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 8, no. 1, pp. 1–38, 2024.
- [36] D. Rivkin, F. Hogan, A. Feriani, A. Konar, A. Sigal, S. Liu, and G. Dudek, "Sage: Smart home agent with grounded execution," *arXiv preprint arXiv:2311.00772*, 2023.
- [37] S. Kerman Freire, M. Foosherian, C. Wang, and E. Niforatos, "Harnessing Large Language Models for cognitive assistants in factories," in *Proceedings of the 5th International Conference on Conversational User Interfaces*, 2023, pp. 1–6.
- [38] B. K. Saha, "Generative Artificial Intelligence for industry: Opportunities, challenges, and impact," in *2024 International Conference on Artificial Intelligence in Information and Communication (ICAIC)*. IEEE, 2024, pp. 081–086.
- [39] İ. Kök, O. Demirci, and S. Özdemir, "When IoT meet LLMs: Applications and challenges," in *2024 IEEE International Conference on Big Data (BigData)*. IEEE, 2024, pp. 7075–7084.
- [40] M. R. Tabassum, M. J. Ritchie, S. Mustafiz, and J. Kienzle, "Using LLMs for use case modelling of IoT systems: An experience report," in *Proceedings of the ACM/IEEE International Conference on Model Driven Engineering Languages and Systems*, 2024, pp. 611–619.
- [41] Y. Chang, X. Wang, J. Wang, Y. Wu, L. Yang, K. Zhu, H. Chen, X. Yi, C. Wang, Y. Wang, W. Ye, Y. Zhang, Y. Chang, P. S. Yu, Q. Yang, and X. Xie, "A survey on evaluation of Large Language Models," *ACM Transactions on Intelligent Systems and Technology*, vol. 15, no. 3, pp. 1–45, 2024.
- [42] T. Cui, Y. Wang, C. Fu, Y. Xiao, S. Li, X. Deng, Y. Liu, Q. Zhang, Z. Qiu, P. Li, Z. Tan, J. Xiong, X. Kong, Z. Wen, K. Xu, and Q. Li, "Risk taxonomy, mitigation, and assessment benchmarks of Large Language Model systems," *arXiv preprint arXiv:2401.05778*, 2024.
- [43] B. C. Das, M. H. Amini, and Y. Wu, "Security and privacy challenges of Large Language Models: A survey," *ACM Computing Surveys*, vol. 57, no. 6, pp. 1–39, 2025.
- [44] O. Friha, M. A. Ferrag, B. Kantarci, B. Cakmak, A. Ozgun, and N. Ghoualmi-Zine, "LLM-based edge intelligence: A comprehensive survey on architectures, applications, security and trustworthiness," *IEEE Open Journal of the Communications Society*, 2024.
- [45] Z. Guo, A. Lai, J. H. Thygesen, J. Farrington, T. Keen, and K. Li, "Large Language Model for mental health: A systematic review," *arXiv preprint arXiv:2403.15401*, 2024.
- [46] S. Hu, T. Huang, F. Ilhan, S. Tekin, G. Liu, R. Kompella, and L. Liu, "A survey on Large Language Model-based game agents," *arXiv preprint arXiv:2404.02039*, 2024.
- [47] M. Jin, Q. Wen, Y. Liang, C. Zhang, S. Xue, X. Wang, J. Zhang, Y. Wang, H. Chen, X. Li, S. Pan, V. S. Tseng, Y. Zheng, L. Chen, and H. Xiong, "Large models for time series and spatio-temporal data: A survey and outlook," *arXiv preprint arXiv:2310.10196*, 2023.
- [48] H. R. Lawrence, R. A. Schneider, S. B. Rubin, M. J. Mataric, D. J. McDuff, and M. J. Bell, "The opportunities and risks of Large Language Models in mental health," *arXiv preprint arXiv:2403.14814*, 2024.
- [49] G. Qu, Q. Chen, W. Wei, Z. Lin, X. Chen, and K. Huang, "Mobile edge intelligence for Large Language Models: A contemporary survey," *IEEE Communications Surveys & Tutorials*, 2025.
- [50] S. Wang, T. Zhu, B. Liu, D. Ming, Z. Guo, D. Ye, and W. Zhou, "Unique security and privacy threats of Large Language Model: A comprehensive survey," *arXiv preprint arXiv:2406.07973*, 2024.
- [51] H. Zhou, C. Hu, Y. Yuan, Y. Cui, Y. Jin, C. Chen, H. Wu, D. Yuan, L. Jiang, D. Wu, X. Liu, J. Zhang, X. Wang, and J. Liu, "Large Language Model (LLM) for telecommunications: A comprehensive survey on principles, key techniques, and opportunities," *IEEE Communications Surveys & Tutorials*, 2024.
- [52] Y. Zhu, H. Yuan, S. Wang, J. Liu, W. Liu, C. Deng, Z. Dou, and J.-R. Wen, "Large Language Models for information retrieval: A survey," *arXiv preprint arXiv:2308.07107*, 2023.

- [53] Z. Yang, Z. Dai, Y. Yang, J. Carbonell, R. R. Salakhutdinov, and Q. V. Le, "Xlnet: Generalized autoregressive pretraining for language understanding," *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [54] A. Zeng, X. Liu, Z. Du, Z. Wang, H. Lai, M. Ding, Z. Yang, Y. Xu, W. Zheng, X. Xia, W. L. Tam, Z. Ma, Y. Xue, J. Zhai, W. Chen, P. Zhang, Y. Dong, and J. Tang, "GLM-130B: An open bilingual pre-trained model," *arXiv preprint arXiv:2210.02414*, 2022.
- [55] A. Chowdhery, S. Narang, J. Devlin, M. Bosma, G. Mishra, A. Roberts, P. Barham, H. W. Chung, C. Sutton, S. Gehrmann, P. Schuh, K. Shi, S. Tsvyashchenko, J. Maynez, A. Rao, P. Barnes, Y. Tay, N. Shazeer, V. Prabhakaran, E. Reif, N. Du, B. Hutchinson, R. Pope, J. Bradbury, J. Austin, M. Isard, G. Gur-Ari, P. Yin, T. Duke, A. Levskaya, S. Ghemawat, S. Dev, H. Michalewski, X. Garcia, V. Misra, K. Robinson, L. Fedus, D. Zhou, D. Ippolito, D. Luan, H. Lim, B. Zoph, A. Spiridonov, R. Sepassi, D. Dohan, S. Agrawal, M. Omernick, A. M. Dai, T. S. Pillai, M. Pellat, A. Lewkowycz, E. Moreira, R. Child, O. Polozov, K. Lee, Z. Zhou, X. Wang, B. Saeta, M. Diaz, O. Firat, M. Catasta, J. Wei, K. Meier-Hellstern, D. Eck, J. Dean, S. Petrov, and N. Fiedel, "Palm: Scaling language modeling with pathways," *Journal of Machine Learning Research*, vol. 24, no. 240, pp. 1–113, 2023.
- [56] Anthropic. (2024) Introducing the next generation of Claude. Accessed on 21 February 2025. [Online]. Available: <https://www.anthropic.com/news/claude-3-family>
- [57] A. Liu, B. Feng, B. Xue, B. Wang, B. Wu, C. Lu, C. Zhao, C. Deng, C. Zhang, C. Ruan *et al.*, "Deepseek-v3 technical report," *arXiv preprint arXiv:2412.19437*, 2024.
- [58] G. Team, P. Georgiev, V. I. Lei, R. Burnell, L. Bai, A. Gulati, G. Tanzer, D. Vincent, Z. Pan, S. Wang *et al.*, "Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context," *arXiv preprint arXiv:2403.05530*, 2024.
- [59] xAI. (2025) Grok 3 beta — the age of reasoning agents. Accessed on 20 February 2025. [Online]. Available: <https://x.ai/blog/grok-3>
- [60] A. Analysis, "Artificial analysis leaderboards," <https://artificialanalysis.ai/leaderboards/models>, accessed: 2025-08-13.
- [61] Y. Zhang, K. Ayush, S. Qiao, A. A. Heydari, G. Narayanswamy, M. A. Xu, A. A. Metwally, S. Xu, J. Garrison, X. Xu, T. Althoff, Y. Liu, P. Kohli, J. Zhan, M. Malhotra, S. Patel, C. Mascolo, X. Liu, D. McDuff, and Y. Yang, "SensorLM: Learning the language of wearable sensors," *arXiv preprint arXiv:2506.09108*, 2025.
- [62] SymphonyAI, "Symphonyai," 2025, accessed: 2025-08-11. [Online]. Available: <https://www.symphonyai.com/>
- [63] A. Candel, J. McKinney, P. Singer, P. Pfeiffer, M. Jeblick, C. M. Lee, and M. V. Conde, "H2O open ecosystem for state-of-the-art Large Language Models," 2023. [Online]. Available: <https://arxiv.org/abs/2310.13012>
- [64] H2O.ai, "H2o.ai," 2025, accessed: 2025-08-11. [Online]. Available: <https://h2o.ai/>
- [65] World Labs. (2025) Marble world model. Accessed: 2026-01-13. [Online]. Available: <https://www.worldlabs.ai/blog/marble-world-model>
- [66] Microsoft Research, "Towards industrial foundation models: Integrating large language models with industrial data intelligence," 2024, accessed: 2025-08-11. [Online]. Available: <https://www.microsoft.com/en-us/research/articles/towards-industrial-foundation-models-integrating-large-language-models-with-industrial-data-intelligence>
- [67] K. S. Jones, "What is the role of NLP in text retrieval?" in *Natural Language Information Retrieval*. Springer, 1999, pp. 1–24.
- [68] Y. Liu, H. He, T. Han, X. Zhang, M. Liu, J. Tian, Y. Zhang, J. Wang, X. Gao, T. Zhong, Y. Pan, S. Xue, Z. Wu, Z. Li, X. Zhang, S. Zhang, X. Hu, T. Zhang, N. Qiang, T. Li, and B. Ge, "Understanding LLMs: A comprehensive overview from training to inference," *arXiv preprint arXiv:2401.02038*, 2024.
- [69] J. Li, T. Tang, W. X. Zhao, J.-Y. Nie, and J.-R. Wen, "Pre-trained language models for text generation: A survey," *ACM Computing Surveys*, vol. 56, no. 9, pp. 1–39, 2024.
- [70] L. Wang, F. Yang, C. Zhang, J. Lu, J. Qian, S. He, P. Zhao, B. Qiao, R. Huang, S. Qin, Q. Su, J. Ye, Y. Zhang, J.-G. Lou, Q. Lin, S. Rajmohan, D. Zhang, and Q. Zhang, "Large Action Models: From Inception to Implementation," *arXiv preprint arXiv:2412.10047*, 2024.
- [71] Y. Chen, Q. Yang, S. He, Z. Shi, J. Chen, and M. Guizani, "FTPipeHD: A fault-tolerant pipeline-parallel distributed training approach for heterogeneous edge devices," *IEEE Transactions on Mobile Computing*, vol. 23, no. 4, pp. 3200–3212, 2023.
- [72] R. Pope, S. Douglas, A. Chowdhery, J. Devlin, J. Bradbury, J. Heek, K. Xiao, S. Agrawal, and J. Dean, "Efficiently scaling transformer inference," *Proceedings of Machine Learning and Systems*, vol. 5, pp. 606–624, 2023.
- [73] S. Sukhbaatar, E. Grave, P. Bojanowski, and A. Joulin, "Adaptive attention span in transformers," *arXiv preprint arXiv:1905.07799*, 2019.
- [74] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. Gonzalez, H. Zhang, and I. Stoica, "Efficient memory management for Large Language Model serving with pagedattention," in *Proceedings of the 29th Symposium on Operating Systems Principles*, 2023, pp. 611–626.
- [75] R. Y. Aminabadi, S. Rajbhandari, A. A. Awan, C. Li, D. Li, E. Zheng, O. Ruwase, S. Smith, M. Zhang, J. Rasley, and Y. He, "DeepSpeed-Inference: Enabling efficient inference of transformer models at unprecedented scale," in *International Conference for High Performance Computing, Networking, Storage and Analysis*. IEEE, 2022, pp. 1–15.
- [76] S. Ge, Y. Zhang, L. Liu, M. Zhang, J. Han, and J. Gao, "Model tells you what to discard: Adaptive KV cache compression for LLMs," *arXiv preprint arXiv:2310.01801*, 2023.
- [77] Z. Lan, M. Chen, S. Goodman, K. Gimpel, P. Sharma, and R. Soricut, "ALBERT: A lite BERT for self-supervised learning of language representations," *arXiv preprint arXiv:1909.11942*, 2019.
- [78] G. Soliman, H. Zaki, and M. Lotfy, "A comparative analysis of encoder only and decoder only models for challenging LLM-generated STEM MCQs using a self-evaluation approach," *Natural Language Processing Journal*, p. 100131, 2025.
- [79] S. Carreira, T. Marques, J. Ribeiro, and C. Grilo, "Revolutionizing mobile interaction: Enabling a 3 billion parameter GPT LLM on mobile," *arXiv preprint arXiv:2310.01434*, 2023.
- [80] W. Yin, M. Xu, Y. Li, and X. Liu, "LLM as a system service on mobile devices," *arXiv preprint arXiv:2403.11805*, 2024.
- [81] J. Stojkovic, E. Chouksey, C. Zhang, I. Gouri, and J. Torrellas, "Towards greener LLMs: Bringing energy-efficiency to the forefront of LLM inference," *arXiv preprint arXiv:2403.20306*, 2024.
- [82] R. Agrawal, H. Kumar, and S. R. Lnu, "Efficient LLMs for edge devices: Pruning, quantization, and distillation techniques," in *2025 International Conference on Machine Learning and Autonomous Systems (ICMLAS)*. IEEE, 2025, pp. 1413–1418.
- [83] L. Deng, G. Li, S. Han, L. Shi, and Y. Xie, "Model compression and hardware acceleration for neural networks: A comprehensive survey," *Proceedings of the IEEE*, vol. 108, no. 4, pp. 485–532, 2020.
- [84] H. Cheng, M. Zhang, and J. Q. Shi, "A survey on deep neural network pruning: Taxonomy, comparison, analysis, and recommendations," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
- [85] W. Zhao, L. Zou, Z. Wang, X. Yao, and B. Yu, "HAPE: Hardware-aware LLM pruning for efficient on-device inference optimization," *ACM Transactions on Design Automation of Electronic Systems*, 2025.
- [86] Z. Yu, Z. Wang, Y. Li, R. Gao, X. Zhou, S. R. Bommur, Y. Zhao, and Y. Lin, "Edge-LLM: Enabling efficient large language model adaptation on edge devices via unified compression and adaptive layer voting," in *ACM/IEEE Design Automation Conference*, 2024, pp. 1–6.
- [87] S. Yao, Y. Zhao, A. Zhang, L. Su, and T. Abdelzaher, "DeepIoT: Compressing deep neural network structures for sensing systems with a compressor-critic framework," in *Proceedings of the 15th ACM Conference on Embedded Networked Sensor Systems*, ser. SenSys '17, 2017. [Online]. Available: <https://doi.org/10.1145/3131672.3131675>
- [88] S. Yao, Y. Zhao, H. Shao, S. Liu, D. Liu, L. Su, and T. Abdelzaher, "FastDeepIoT: Towards understanding and optimizing neural network execution time on mobile and embedded devices," in *Proceedings of the International Conference on Mobile Computing and Networking*. ACM, 2018, pp. 1–14.
- [89] X. Hu, Y. Cheng, D. Yang, Z. Yuan, J. Yu, C. Xu, and S. Zhou, "I-LLM: Efficient integer-only inference for fully-quantized low-bit large language models," *arXiv preprint arXiv:2405.17849*, 2024.
- [90] J. Kong, L. Hu, F. Ponzina, and T. Rosing, "TinyAgent: Quantization-aware model compression and adaptation for on-device LLM agent deployment," in *Workshop on Efficient Systems for Foundation*

- Models*, 2024. [Online]. Available: <https://openreview.net/forum?id=ntl0eq0urB>
- [91] P. Micikevicius, S. Narang, J. Alben, G. Diamos, E. Elsen, D. Garcia, B. Ginsburg, M. Houston, O. Kuchaiev, G. Venkatesh, and H. Wu, "Mixed precision training," *arXiv preprint arXiv:1710.03740*, 2017.
 - [92] C. Yang, Y. Zhu, W. Lu, Y. Wang, Q. Chen, C. Gao, B. Yan, and Y. Chen, "Survey on knowledge distillation for Large Language Models: methods, evaluation, and application," *ACM Transactions on Intelligent Systems and Technology*, 2024.
 - [93] R. Jin, Q. Xu, M. Wu, Y. Xu, D. Li, X. Li, and Z. Chen, "LLM-based knowledge pruning for time series data analytics on edge-computing devices," *arXiv preprint arXiv:2406.08765*, 2024.
 - [94] A. Kundu, F. Lim, A. Chew, L. Wynter, P. Chong, and R. D. Lee, "Efficiently distilling LLMs for edge applications," *arXiv preprint arXiv:2404.01353*, 2024.
 - [95] X. Bi, D. Chen, G. Chen, S. Chen, D. Dai, C. Deng, H. Ding, K. Dong, Q. Du, Z. Fu, H. Gao, K. Gao, W. Gao, R. Ge, K. Guan, D. Guo, J. Guo, G. Hao, Z. Hao, Y. He, W. Hu, P. Huang, E. Li, G. Li, J. Li, Y. Li, Y. Li, W. Liang, F. Lin, A. Liu, B. Liu, W. Liu, X. Liu, X. Liu, Y. Liu, H. Lu, S. Lu, F. Luo, S. Ma, X. Nie, T. Pei, Y. Piao, J. Qiu, H. Qu, T. Ren, Z. Ren, C. Ruan, Z. Sha, Z. Shao, J. Song, X. Su, J. Sun, Y. Sun, M. Tang, B. Wang, P. Wang, S. Wang, Y. Wang, Y. Wang, T. Wu, Y. Wu, X. Xie, Z. Xie, Z. Xie, Y. Xiong, H. Xu, R. Xu, Y. Xu, D. Yang, Y. You, S. Yu, X. Yu, B. Zhang, H. Zhang, L. Zhang, L. Zhang, M. Zhang, M. Zhang, W. Zhang, Y. Zhang, C. Zhao, Y. Zhao, S. Zhou, S. Zhou, Q. Zhu, and Y. Zou, "DeepSeek LLM: Scaling open-source language models with longtermism," *arXiv e-prints*, pp. arXiv-2401, 2024.
 - [96] M. U. Nasir, S. Earle, J. Togelius, S. James, and C. Cleghorn, "LLMatic: neural architecture search via large language models and quality diversity optimization," in *Proceedings of the Genetic and Evolutionary Computation Conference*, 2024, pp. 1110–1118.
 - [97] J. Yue, J. Zhao, and C. Zhao, "Less but better: Enabling generalized zero-shot learning towards unseen domains by intrinsic learning from redundant LLM semantics," *arXiv preprint arXiv:2403.14362*, 2024.
 - [98] C. H. Lampert, H. Nickisch, and S. Harmeling, "Attribute-based classification for zero-shot visual object categorization," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 36, no. 3, pp. 453–465, 2013.
 - [99] S. Ji, X. Zheng, and C. Wu, "HARGPT: Are LLMs zero-shot human activity recognizers?" *arXiv preprint arXiv:2403.02727*, 2024.
 - [100] Y. Kim, X. Xu, D. McDuff, C. Breazeal, and H. W. Park, "Health-LLM: Large Language Models for health prediction via wearable sensor data," *arXiv preprint arXiv:2401.06866*, 2024.
 - [101] T. Okita, K. Ukita, K. Matsuishi, M. Kagiya, K. Hirata, and A. Miyazaki, "Towards LLMs for sensor data: Multi-task self-supervised learning," in *ACM International Joint Conference on Pervasive and Ubiquitous Computing & the ACM International Symposium on Wearable Computing*, 2023, pp. 499–504.
 - [102] X. Zeng, X. Wang, T. Zhang, C. Yu, S. Zhao, and Y. Chen, "GestureGPT: Zero-shot interactive gesture understanding and grounding with large language model agents," *arXiv preprint arXiv:2310.12821*, 2023.
 - [103] X. Xu, B. Yao, Y. Dong, S. Gabriel, H. Yu, J. Hendler, M. Ghassemi, A. K. Dey, and D. Wang, "Mental-LLM: Leveraging Large Language Models for mental health prediction via online text data," *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 8, no. 1, pp. 1–32, 2024.
 - [104] B. Chen, Z. Zhang, N. Langrené, and S. Zhu, "Unleashing the potential of prompt engineering in Large Language Models: A comprehensive review," *arXiv preprint arXiv:2310.14735*, 2023.
 - [105] J. Gao, Y. Zhang, Y. Chen, T. Zhang, B. Tang, and X. Wang, "Unsupervised human activity recognition via Large Language Models and iterative evolution," in *IEEE International Conference on Acoustics, Speech and Signal Processing*. IEEE, 2024, pp. 91–95.
 - [106] Z. Hu, Y. Zhang, R. Rossi, T. Yu, S. Kim, and S. Pan, "Are Large Language Models capable of causal reasoning for sensing data analysis?" in *Proceedings of the Workshop on Edge and Mobile Foundation Models*, 2024, pp. 24–29.
 - [107] T. Han, Z. Wang, C. Fang, S. Zhao, S. Ma, and Z. Chen, "Token-budget-aware llm reasoning," in *Findings of the Association for Computational Linguistics: ACL 2025*, 2025, pp. 24 842–24 855.
 - [108] Z. Li, Q. Dong, J. Ma, D. Zhang, K. Jia, and Z. Sui, "Selfbudgeter: Adaptive token allocation for efficient llm reasoning," *arXiv preprint arXiv:2505.11274*, 2025.
 - [109] T. Liu, Z. Wang, J. Miao, I. Hsu, J. Yan, J. Chen, R. Han, F. Xu, Y. Chen, K. Jiang et al., "Budget-aware tool-use enables effective agent scaling," *arXiv preprint arXiv:2511.17006*, 2025.
 - [110] R. Sapkota, K. I. Roumeliotis, and M. Karkee, "Ai agents vs. agentic ai: A conceptual taxonomy, applications and challenges," *arXiv preprint arXiv:2505.10468*, 2025.
 - [111] P. Belcak, G. Heinrich, S. Diao, Y. Fu, X. Dong, S. Muralidharan, Y. C. Lin, and P. Molchanov, "Small language models are the future of agentic ai," *arXiv preprint arXiv:2506.02153*, 2025.
 - [112] M. A. Farahani, M. I. Khan, and T. Wuest, "Hybrid agentic ai and multi-agent systems in smart manufacturing," *arXiv preprint arXiv:2511.18258*, 2025.
 - [113] Y. Du, Q. Yang, L. Wang, J. Lin, H. Cui, and S. C. Liew, "Llmind 2.0: Distributed iot automation with natural language m2m communication and lightweight llm agents," *arXiv preprint arXiv:2508.13920*, 2025.
 - [114] Q. Guan, J. Ouyang, D. Wu, and W. Yu, "Citygpt: Towards urban iot learning, analysis and interaction with multi-agent system," *arXiv preprint arXiv:2405.14691*, 2024.
 - [115] A. Nedisan Videsjorden, H. Song, A. Goknil, D. Roman, and A. Soylu, "Lumen: Enhancing iot system observability with multi-agent large language models and knowledge graphs," *ACM Transactions on Internet of Things*, 2025.
 - [116] R. J. Mathav, K. VM, H. Warriar, and Y. Gupta, "Fine tuning LLM for enterprise: Practical guidelines and recommendations," *arXiv preprint arXiv:2404.10779*, 2024.
 - [117] Y. Wang, Y. Kordi, S. Mishra, A. Liu, N. A. Smith, D. Khashabi, and H. Hajishirzi, "Self-instruct: Aligning language models with self-generated instructions," *arXiv preprint arXiv:2212.10560*, 2022.
 - [118] R. He, L. Liu, H. Ye, Q. Tan, B. Ding, L. Cheng, J.-W. Low, L. Bing, and L. Si, "On the effectiveness of adapter-based tuning for pretrained language model adaptation," *arXiv preprint arXiv:2106.03164*, 2021.
 - [119] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRa: Low-rank adaptation of Large Language Models," *arXiv preprint arXiv:2106.09685*, 2021.
 - [120] S. Liu, J. Keung, Z. Yang, F. Liu, Q. Zhou, and Y. Liao, "Delving into parameter-efficient fine-tuning in code change learning: An empirical study," *arXiv preprint arXiv:2402.06247*, 2024.
 - [121] X. Chen and X. Zhang, "RF Genesis: Zero-shot generalization of mmwave sensing through simulation-based data synthesis and generative diffusion models," in *Proceedings of the 21st ACM Conference on Embedded Networked Sensor Systems*, 2023, pp. 28–42.
 - [122] J. Yuan, C. Yang, D. Cai, S. Wang, X. Yuan, Z. Zhang, X. Li, D. Zhang, H. Mei, X. Jia, S. Wang, and M. Xu, "Mobile foundation model as firmware," in *Proceedings of the Annual International Conference on Mobile Computing and Networking*, 2024, pp. 279–295.
 - [123] Y. Ren, Y. Chen, S. Liu, B. Wang, H. Yu, and Z. Cui, "TPLLM: A traffic prediction framework based on pretrained Large Language Models," *arXiv preprint arXiv:2403.02221*, 2024.
 - [124] F. Yin, X. Ye, and G. Durrett, "LoFiT: Localized fine-tuning on LLM representations," *arXiv preprint arXiv:2406.01563*, 2024.
 - [125] N. Muennighoff, Z. Yang, W. Shi, X. L. Li, L. Fei-Fei, H. Hajishirzi, L. Zettlemoyer, P. Liang, E. Candès, and T. Hashimoto, "s1: Simple test-time scaling," *arXiv preprint arXiv:2501.19393*, 2025.
 - [126] J. Geiping, S. McLeish, N. Jain, J. Kirchenbauer, S. Singh, B. R. Bartoldson, B. Kailkhura, A. Bhatele, and T. Goldstein, "Scaling up test-time compute with latent reasoning: A recurrent depth approach," *arXiv preprint arXiv:2502.05171*, 2025.
 - [127] S. Oh, S. Kum, and J. Moon, "Real-time environment monitoring and response through IoT and retrieval-augmented generation," in *International Conference on Information and Communication Technology Convergence*. IEEE, 2024, pp. 1658–1659.
 - [128] T. Ouyang, G. Hong, K. Zhao, Z. Zhou, W. Wu, Z. Lv, and X. Chen, "Adarag: Adaptive optimization for retrieval augmented generation with multilevel retrievers at the edge," in *IEEE INFOCOM 2025-IEEE Conference on Computer Communications*. IEEE, 2025, pp. 1–10.
 - [129] C. Su, J. Wen, J. Kang, Y. Wang, Y. Su, H. Pan, Z. Zhong, and M. S. Hossain, "Hybrid RAG-empowered multi-modal LLM for secure data

- management in internet of medical things: A diffusion-based contract approach," *IEEE Internet of Things Journal*, 2024.
- [130] Y. Dong, Y. L. Aung, S. Chattopadhyay, and J. Zhou, "ChatIoT: Large language model-based security assistant for Internet of Things with retrieval-augmented generation," *arXiv preprint arXiv:2502.09896*, 2025.
- [131] T. Park, G. Lee, and M.-S. Kim, "MobileRAG: A fast, memory-efficient, and energy-efficient method for on-device RAG," *arXiv preprint arXiv:2507.01079*, 2025.
- [132] B. Reichman, X. Yu, L. Hu, J. Truxal, A. Jain, R. Chandrupatla, T. Š. Rosing, and L. Heck, "SensorQA: A question answering benchmark for daily-life monitoring," *arXiv preprint arXiv:2501.04974*, 2025.
- [133] H. Huang, Y. Qu, J. Liu, M. Yang, and T. Zhao, "An empirical study of LLM-as-a-judge for LLM evaluation: Fine-tuned judge models are task-specific classifiers," *arXiv preprint arXiv:2403.02839*, 2024.
- [134] M. A. Haque, F. Rahman, K. D. Gupta, K. Shujaee, and R. George, "Tinyllm: Evaluation and optimization of small language models for agentic tasks on edge devices," *arXiv preprint arXiv:2511.22138*, 2025.
- [135] Y. Zhang, L. Cui, W. Bi, and S. Shi, "Alleviating hallucinations of Large Language Models through induced hallucinations," *arXiv preprint arXiv:2312.15710*, 2023.
- [136] L. Bürger, F. A. Hamprecht, and B. Nadler, "Truth is universal: Robust detection of lies in LLMs," *arXiv preprint arXiv:2407.12831*, 2024.
- [137] F. Zhang, P. Yu, B. Yi, B. Zhang, T. Li, and Z. Liu, "Prompt-guided internal states for hallucination detection of Large Language Models," *arXiv preprint arXiv:2411.04847*, 2024.
- [138] S. Marks and M. Tegmark, "The geometry of truth: Emergent linear structure in large language model representations of true/false datasets," *arXiv preprint arXiv:2310.06824*, 2023.
- [139] K. Verspoor, "Fighting fire with fire—using LLMs to combat LLM hallucinations," 2024.
- [140] J. Zhang, Z. Li, K. Das, B. A. Malin, and S. Kumar, "SAC3: Reliable hallucination detection in black-box language models via semantic-aware cross-check consistency," *arXiv preprint arXiv:2311.01740*, 2023.
- [141] A.-R. Ottun, R. Marasinghe, T. Elemosho, M. Liyanage, M. Ragab, P. Bagave, M. Westberg, M. Asadi, M. Boerger, C. Sandeepa, T. Senevirathna, B. Siniarski, M. Liyanage, V. H. La, M.-D. Nguyen, E. Montes de Oca, T. Oomen, J. F. F. Goncalves, I. Tanaskovic, S. Klopovic, N. Kourtellis, C. Soriente, J. Pridmore, A. R. Cavalli, D. Draskovic, S. Marchal, S. Wang, D. Solans Noguero, N. Tcholtchev, A. Y. Ding, and H. Flores, "The SPATIAL architecture: Design and development experiences from gauging and monitoring the AI inference capabilities of modern applications," in *2024 IEEE 44th International Conference on Distributed Computing Systems (ICDCS)*. IEEE, 2024, pp. 947–959.
- [142] Y. Liu, Y. Yao, J.-F. Ton, X. Zhang, R. G. H. Cheng, Y. Klochkov, M. F. Taufiq, and H. Li, "Trustworthy LLMs: A survey and guideline for evaluating Large Language Models' alignment," *arXiv preprint arXiv:2308.05374*, 2023.
- [143] L. Sun, Y. Huang, H. Wang, S. Wu, Q. Zhang, Y. Li, C. Gao, Y. Huang, W. Lyu, Y. Zhang, X. Li, H. Sun, Z. Liu, Y. Liu, Y. Wang, Z. Zhang, B. Vidgen, B. Kailkhura, C. Xiong, C. Xiao, C. Li, E. Xing, F. Huang, H. Liu, H. Ji, H. Wang, H. Zhang, H. Yao, M. Kellis, M. Zitnik, M. Jiang, M. Bansal, J. Zou, J. Pei, J. Liu, J. Gao, J. Han, J. Zhao, J. Tang, J. Wang, J. Vanschoren, J. Mitchell, K. Shu, K. Xu, K.-W. Chang, L. He, L. Huang, M. Backes, N. Z. Gong, P. S. Yu, P.-Y. Chen, Q. Gu, R. Xu, R. Ying, S. Ji, S. Jana, T. Chen, T. Liu, T. Zhou, W. Wang, X. Li, X. Zhang, X. Wang, X. Xie, X. Chen, X. Wang, Y. Liu, Y. Ye, Y. Cao, Y. Chen, and Y. Zhao, "TrustLLM: Trustworthiness in Large Language Models," *arXiv preprint arXiv:2401.05561*, 2024.
- [144] M. A. Ferrag, M. Ndhlovu, N. Tihanyi, L. C. Cordeiro, M. Debbah, T. Lestable, and N. S. Thandi, "Revolutionizing cyber threat detection with Large Language Models: A privacy-preserving Bert-based lightweight model for IoT/IIoT devices," *IEEE Access*, 2024.
- [145] Z. Zhang, M. Jia, H.-P. Lee, B. Yao, S. Das, A. Lerner, D. Wang, and T. Li, "It's a fair game, or is it? Examining how users navigate disclosure risks and benefits when using LLM-based conversational agents," in *Proceedings of the CHI Conference on Human Factors in Computing Systems*. ACM, 2024, pp. 156–1.
- [146] X. Wu, R. Duan, and J. Ni, "Unveiling security, privacy, and ethical concerns of ChatGPT," *Journal of Information and Intelligence*, vol. 2, no. 2, pp. 102–115, 2024.
- [147] D. Bowen, B. Murphy, W. Cai, D. Khachaturov, A. Gleave, and K. Pelrine, "Data poisoning in LLMs: Jailbreak-tuning and scaling laws," *arXiv preprint arXiv:2408.02946*, 2024.
- [148] M. Miranda, E. S. Ruzzetti, A. Santilli, F. M. Zanzotto, S. Bratières, and E. Rodolà, "Preserving privacy in Large Language Models: A survey on current threats and solutions," *arXiv preprint arXiv:2408.05212*, 2024.
- [149] H. Li, Y. Chen, J. Luo, J. Wang, H. Peng, Y. Kang, X. Zhang, Q. Hu, C. Chan, Z. Xu, B. Hooi, and Y. Song, "Privacy in Large Language Models: Attacks, defenses and future directions," *arXiv preprint arXiv:2310.10383*, 2023.
- [150] S. Han, B. Buyukates, Z. Hu, H. Jin, W. Jin, L. Sun, X. Wang, W. Wu, C. Xie, Y. Yao, K. Zhang, Q. Zhang, Y. Zhang, C. Joe-Wong, S. Avestimehr, and C. He, "FedSecurity: A benchmark for attacks and defenses in federated learning and federated LLMs," in *Proceedings of the ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2024, pp. 5070–5081.
- [151] J. Zheng, H. Zhang, L. Wang, W. Qiu, H. Zheng, and Z. Zheng, "Safely learning with private data: A federated learning framework for Large Language Model," *arXiv preprint arXiv:2406.14898*, 2024.
- [152] L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin, and T. Liu, "A survey on hallucination in Large Language Models: Principles, taxonomy, challenges, and open questions," *ACM Transactions on Information Systems*, 2024.
- [153] J. Jiao, S. Afroogh, Y. Xu, and C. Phillips, "Navigating LLM ethics: Advancements, challenges, and future directions," *arXiv preprint arXiv:2406.18841*, 2024.
- [154] P. Yubeaton, T. Mahmoud, S. Naga, P. Taheri, T. Xia, A. George, Y. Khalil, S. Q. Zhang, S. Joshi, C. Hegde, and S. Garg, "Huff-LLM: End-to-end lossless compression for efficient LLM inference," *arXiv preprint arXiv:2502.00922*, 2025.
- [155] Y. Zhu, X. Wang, J. Chen, S. Qiao, Y. Ou, Y. Yao, S. Deng, H. Chen, and N. Zhang, "LLMs for knowledge graph construction and reasoning: Recent capabilities and future opportunities," *World Wide Web*, vol. 27, no. 5, p. 58, 2024.
- [156] C.-Y. Chiu, D.-Y. Hong, P. Liu, and J.-J. Wu, "Effective compression of language models by combining pruning and knowledge distillation," in *IEEE Annual Computers, Software, and Applications Conference*. IEEE, 2024, pp. 429–438.
- [157] W. Wang, F. Wei, L. Dong, H. Bao, N. Yang, and M. Zhou, "MiniLM: Deep self-attention distillation for task-agnostic compression of pre-trained transformers," *Advances in Neural Information Processing Systems*, vol. 33, pp. 5776–5788, 2020.
- [158] P. P. Ray, "A review on llms for iot ecosystem: state-of-the-art, lightweight models, use cases, key challenges, future directions," 2025.
- [159] X. Chen, W. Wu, L. Li, and F. Ji, "Llm-empowered iot for 6g networks: Architecture, challenges, and solutions," *IEEE Internet of Things Magazine*, vol. 8, no. 6, p. 34–41, Nov 2025.
- [160] A. Abouelenin, A. Ashfaq, A. Atkinson, H. Awadalla, N. Bach, J. Bao, A. Benhaim, M. Cai, V. Chaudhary, C. Chen *et al.*, "Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-loras," *arXiv preprint arXiv:2503.01743*, 2025.
- [161] G. Team, A. Kamath, J. Ferret, S. Pathak, N. Vieillard, R. Merhej, S. Perrin, T. Matejovicova, A. Ramé, M. Rivière *et al.*, "Gemma 3 technical report," *arXiv preprint arXiv:2503.19786*, 2025.
- [162] I. Ahmed, S. Islam, P. P. Datta, I. Kabir, N. U. R. Chowdhury, and A. Haque, "Qwen 2.5: A comprehensive review of the leading resource-efficient llm with potential to surpass all competitors," *Authorea Preprints*, 2025.
- [163] W. Zhao, W. Jing, Z. Lu, and X. Wen, "Edge and terminal cooperation enabled LLM deployment optimization in wireless network," in *IEEE/CIC International Conference on Communications in China*. IEEE, 2024, pp. 220–225.
- [164] F. Cai, D. Yuan, Z. Yang, and L. Cui, "Edge-LLM: A collaborative framework for Large Language Model serving in edge computing," in *International Conference on Web Services*. IEEE, 2024, pp. 799–809.
- [165] J. Li, B. Han, S. Li, X. Wang, and J. Li, "CoLLM: A collaborative LLM inference framework for resource-constrained devices," in *In-*

- ternational Conference on Communications in China. IEEE, 2024, pp. 185–190.
- [166] H. Wang, Q. Zhou, Z. Hong, and S. Guo, “D² MoE: Dual Routing and Dynamic Scheduling for Efficient On-Device MoE-based LLM Serving,” *arXiv preprint arXiv:2504.15299*, 2025.
- [167] N. Li, S. Guo, T. Zhang, M. Li, Z. Hong, Q. Zhou, X. Yuan, and H. Zhang, “The MoE-Empowered Edge LLMs Deployment: Architecture, challenges, and opportunities,” *arXiv preprint arXiv:2502.08381*, 2025.
- [168] L. Jin, Y. Zhang, Y. Li, S. Wang, H. H. Yang, J. Wu, and M. Zhang, “MoE²: Optimizing collaborative inference for edge large language models,” *arXiv preprint arXiv:2501.09410*, 2025.
- [169] E. Malach, G. Yehudai, S. Shalev-Schwartz, and O. Shamir, “Proving the lottery ticket hypothesis: Pruning is all you need,” in *International Conference on Machine Learning*. PMLR, 2020, pp. 6682–6691.
- [170] A. Li, J. Sun, B. Wang, L. Duan, S. Li, Y. Chen, and H. Li, “LotteryFL: Personalized and communication-efficient federated learning with lottery ticket hypothesis on non-iid datasets,” *arXiv preprint arXiv:2008.03371*, 2020.
- [171] J. Wang, J. Wang, B. Athiwaratkun, C. Zhang, and J. Zou, “Mixture-of-agents enhances large language model capabilities,” *arXiv preprint arXiv:2406.04692*, 2024.
- [172] J. Li, J. Xu, S. Huang, Y. Chen, W. Li, J. Liu, Y. Lian, J. Pan, L. Ding, H. Zhou, Y. Wang, and G. Dai, “Large language model inference acceleration: A comprehensive hardware perspective,” *arXiv preprint arXiv:2410.04466*, 2024.
- [173] M. Alqarni and A. Azim, “SecureLLAMA: Secure FPGAs using LLAMA Large Language Models,” *IEEE Transactions on Artificial Intelligence*, 2025.
- [174] J. Haris, R. Saha, W. Hu, and J. Cano, “Designing efficient llm accelerators for edge devices,” *arXiv preprint arXiv:2408.00462*, 2024.
- [175] M. Huang, A. Shen, K. Li, H. Peng, B. Li, Y. Su, and H. Yu, “EdgeLLM: A highly efficient CPU-FPGA heterogeneous edge accelerator for large language models,” *IEEE Transactions on Circuits and Systems I: Regular Papers*, 2025.
- [176] D. Xu, H. Zhang, L. Yang, R. Liu, G. Huang, M. Xu, and X. Liu, “Fast on-device LLM inference with NPUS,” in *Proceedings of the ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1*, 2025, pp. 445–462.
- [177] L. Chen, Y. Zhao, Q. Xie, and Q. Sheng, “Optimization of Armv9 architecture general large language model inference performance based on Llama. cpp,” *arXiv preprint arXiv:2406.10816*, 2024.
- [178] Z. Jian, X. Liu, Q. Dong, L. Cheng, X. Xie, and T. Li, “Smartzone: Runtime support for secure and efficient on-device inference on arm trustzone,” *IEEE Transactions on Computers*, 2025.
- [179] Y. Gao, K. Xiao, F. Li, W. Xu, J. Huang, and W. Dong, “ChatIoT: Zero-code generation of trigger-action based IoT programs,” *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 8, no. 3, pp. 1–29, 2024.
- [180] L. Shen and Y. Zheng, “Iotcoder: A copilot for IoT application development,” in *Proceedings of the Annual International Conference on Mobile Computing and Networking*, 2024, pp. 1647–1649.
- [181] G. Wu, L. Hu, Y. Hu, X. Xiong, and F. Wang, “LLM4TAP: LLM-enhanced TAP rule recommendation,” *IEEE Internet of Things Journal*, 2025.
- [182] H. Su, “LLM-based code development model with active prompt adjustment and on-device validation,” *Authorea Preprints*, 2025.
- [183] Q. Dong, X. Chen, and M. Satyanarayanan, “Creating edge AI from cloud-based LLMs,” in *Proceedings of the 25th International Workshop on Mobile Computing Systems and Applications*, 2024, pp. 8–13.
- [184] Y. Shen, J. Shao, X. Zhang, Z. Lin, H. Pan, D. Li, J. Zhang, and K. B. Letaief, “Large Language Models empowered autonomous edge AI for connected intelligence,” *IEEE Communications Magazine*, 2024.
- [185] H. Yang, Y. Zhou, T. Liang, and L. Kuang, “ChatDL: An LLM-based defect localization approach for software in IIoT flexible manufacturing,” *IEEE Internet of Things Journal*, 2025.
- [186] Y. Tang, A. A. B. Da Costa, X. Zhang, I. Patrick, S. Khastgir, and P. Jennings, “Domain knowledge distillation from Large Language Model: An empirical study in the autonomous driving domain,” in *IEEE International Conference on Intelligent Transportation Systems*. IEEE, 2023, pp. 3893–3900.
- [187] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, “ProgPrompt: Generating situated robot task plans using Large Language Models,” in *Proceedings of the IEEE International Conference on Robotics and Automation*. IEEE, 2023, pp. 11 523–11 530.
- [188] B. Adnan, S. Miryala, A. Sambu, K. Vaidhyathanan, M. De Sanctis, and R. Spalazzese, “Leveraging llms for dynamic IoT systems generation through mixed-initiative interaction,” *arXiv preprint arXiv:2502.00689*, 2025.
- [189] Y. Liu, L. Palmieri, S. Koch, I. Georgievski, and M. Aiello, “Delta: Decomposed efficient long-term robot task planning using large language models,” *arXiv preprint arXiv:2404.03275*, 2024.
- [190] H. Wei, Z. Zhang, S. He, T. Xia, S. Pan, and F. Liu, “PlanGenLLMs: A modern survey of LLM planning capabilities,” *arXiv preprint arXiv:2502.11221*, 2025.
- [191] J. Hu, H. Jia, M. Hassan, L. Yao, B. Kusy, and W. Hu, “LightLLM: A versatile large language model for predictive light sensing,” in *Proceedings of the 23rd ACM Conference on Embedded Networked Sensor Systems*, 2025, pp. 158–171.
- [192] N. Zhong, Y. Wang, R. Xiong, Y. Zheng, Y. Li, M. Ouyang, D. Shen, and X. Zhu, “CASIT: Collective intelligent agent system for Internet of Things,” *IEEE Internet of Things Journal*, 2024.
- [193] H. Yang, R. Wu, and W. Xu, “TransCompressor: LLM-powered multimodal data compression for smart transportation,” in *Proceedings of the Annual International Conference on Mobile Computing and Networking*, 2024, pp. 2335–2340.
- [194] J. Wen, J. Nie, J. Kang, D. Niyato, H. Du, Y. Zhang, and M. Guizani, “From generative AI to generative Internet of Things: Fundamentals, framework, and outlooks,” *IEEE Internet of Things Magazine*, vol. 7, no. 3, pp. 30–37, 2024.
- [195] S. Tong, H. Liu, R. Guo, X. Tian, W. Wang, D. Liu, and Y. Zhang, “A text-based knowledge-embedded soft sensing modeling approach for general industrial process tasks based on Large Language Model,” *arXiv preprint arXiv:2501.05075*, 2025.
- [196] T. An, Y. Zhou, H. Zou, and J. Yang, “Iot-llm: A framework for enhancing large language model reasoning from real-world sensor data,” Nov 2025. [Online]. Available: <https://arxiv.org/abs/2410.02429>
- [197] A. Soliman, “The future of internet of things and multimodal language models in 6g networks: Opportunities and challenges,” Apr 2025. [Online]. Available: <https://arxiv.org/abs/2504.13971>
- [198] X. Zhou, Y. Hu, Q. Jia, and R. Xie, “Cross-domain integration for general sensor data synthesis: Leveraging LLMs and domain-specific generative models in collaborative environments,” *IEEE Sensors Journal*, 2024.
- [199] J. Zhang, Y. Zhang, X. Cun, Y. Zhang, H. Zhao, H. Lu, X. Shen, and Y. Shan, “Generating human motion from textual descriptions with discrete representations,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023, pp. 14 730–14 740.
- [200] Z. Leng, A. Bhattacharjee, H. Rajasekhar, L. Zhang, E. Bruda, H. Kwon, and T. Plötz, “Imugpt 2.0: Language-based cross modality transfer for sensor-based human activity recognition,” *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 8, no. 3, pp. 1–32, 2024.
- [201] H. Haresamudram, C. I. Tang, S. Suh, P. Lukowicz, and T. Ploetz, “Past, present, and future of sensor-based human activity recognition using wearables: A surveying tutorial on a still challenging task,” *arXiv preprint arXiv:2411.14452*, 2024.
- [202] L. S. S. Ray, B. Zhou, S. Suh, L. Krupp, V. F. Rey, and P. Lukowicz, “Text me the data: Generating ground pressure sequence from textual descriptions for HAR,” in *IEEE International Conference on Pervasive Computing and Communications Workshops and other Affiliated Events (PerCom Workshops)*. IEEE, 2024, pp. 461–464.
- [203] X. Zhou, Q. Jia, and Y. Hu, “Advancing general sensor data synthesis by integrating LLMs and domain-specific generative models,” *IEEE Sensors Letters*, 2024.
- [204] X. Zhou, Q. Jia, Y. Hu, R. Xie, T. Huang, and F. R. Yu, “GenG: An LLM-based generic time series data generation approach for edge intelligence via cross-domain collaboration,” in *IEEE Conference on Computer Communications Workshops*. IEEE, 2024, pp. 1–6.
- [205] S. Paul, L. Zhang, Y. Shen, and H. Jin, “Enabling device control planning capabilities of small language model,” in *IEEE International*

- Conference on Acoustics, Speech and Signal Processing. IEEE, 2024, pp. 12 066–12 070.
- [206] C. Xu, G. Saranathan, M. P. Alam, A. Shah, J. Lim, S. Y. Wong, F. Martin, and S. Bhattacharya, “Data efficient evaluation of Large Language Models and text-to-image models via adaptive sampling,” *arXiv preprint arXiv:2406.15527*, 2024.
 - [207] A. Hota, S. Chatterjee, and S. Chakraborty, “Evaluating Large Language Models as virtual annotators for time-series physical sensing data,” *arXiv preprint arXiv:2403.01133*, 2024.
 - [208] M. Thukral, S. G. Dhekane, S. K. Hiremath, H. Haresamudram, and T. Ploetz, “Layout agnostic human activity recognition in smart homes through textual descriptions of sensor triggers (TDOST),” *arXiv preprint arXiv:2405.12368*, 2024.
 - [209] Z. Huang, P. Peng, F. Lu, and H. Zhang, “An LLM-based method for quality information extraction from web text for crowd-sensing spatiotemporal data,” *Transactions in GIS*, vol. 29, no. 1, 2025.
 - [210] J. Liu, Y. Li, L. Li, Y. Sun, H. Wen, X. Li, Y. Guo, and Y. Liu, “Chainstream: An LLM-based framework for unified synthetic sensing,” *arXiv preprint arXiv:2412.15240*, 2024.
 - [211] Y. Sun and J. Ortiz, “An AI-based system utilizing IoT-enabled ambient sensors and LLMs for complex activity tracking,” *arXiv preprint arXiv:2407.02606*, 2024.
 - [212] H. Guo, X. Su, C. Wu, B. Du, L. Zhang, and D. Li, “Remote sensing ChatGPT : Solving remote sensing tasks with ChatGPT and visual models,” *arXiv preprint arXiv:2401.09083*, 2024.
 - [213] H. Li, M. Xiao, K. Wang, D. I. Kim, and M. Debbah, “Large Language Model based multi-objective optimization for integrated sensing and communications in UAV networks,” *IEEE Wireless Communications Letters*, 2025.
 - [214] W. Zhang, M. Cai, T. Zhang, Y. Zhuang, and X. Mao, “EarthGPT: A universal multi-modal Large Language Model for multi-sensor image comprehension in remote sensing domain,” *IEEE Transactions on Geoscience and Remote Sensing*, 2024.
 - [215] B. Mario and V. Mezhyuev, “A framework for creating an IoT system specification with ChatGPT,” *Internet of Things*, p. 101218, 2024.
 - [216] S. Long, J. Tan, B. Mao, F. Tang, Y. Li, M. Zhao, and N. Kato, “A survey on intelligent network operations and performance optimization based on Large Language Models,” *IEEE Communications Surveys & Tutorials*, 2025.
 - [217] J. Wang, X. Fang, D. Niyato, and T. Liu, “Next-Generation Wi-Fi networks with generative AI: Design and insights,” *arXiv preprint arXiv:2408.04835*, 2024.
 - [218] A. Leivadreas and M. Falkner, “A survey on intent-based networking,” *IEEE Communications Surveys & Tutorials*, vol. 25, no. 1, pp. 625–655, 2022.
 - [219] E. Zeydan and Y. Turk, “Recent advances in intent-based networking: A survey,” in *Vehicular Technology Conference*. IEEE, 2020, pp. 1–5.
 - [220] M. Zhang, X. Shen, J. Cao, Z. Cui, and S. Jiang, “EdgesHard: Efficient LLM inference via collaborative edge computing,” *IEEE Internet of Things Journal*, 2024.
 - [221] C. Gao and S. Q. Zhang, “DL0Ra: Distributed parameter-efficient fine-tuning solution for Large Language Model,” *arXiv preprint arXiv:2404.05182*, 2024.
 - [222] P. Patel, E. Choukse, C. Zhang, Í. Goiri, B. Warriar, N. Mahalingam, and R. Bianchini, “POLCA: Power oversubscription in LLM cloud providers,” *arXiv preprint arXiv:2308.12908*, 2023.
 - [223] A. Sarabi, T. Yin, and M. Liu, “An LLM-based framework for fingerprinting internet-connected devices,” in *Proceedings of the ACM on Internet Measurement Conference*. New York, NY, USA: ACM, 2023, p. 478–484.
 - [224] B. Meyuhas, A. Bremner-Barr, and T. Shapira, “IoT device labeling using Large Language Models,” *arXiv preprint arXiv:2403.01586*, 2024.
 - [225] M. Soleimani, G. Jia, I. Gim, S.-s. Lee, and A. Khandelwal, “Wiretapping llms: Network side-channel attacks on interactive llm services,” *Cryptology ePrint Archive*, 2025.
 - [226] T. Nguyen, H. Nguyen, A. Ijaz, S. Sheikhi, A. V. Vasilakos, and P. Kostakos, “Large Language Models in 6G security: Challenges and opportunities,” *arXiv preprint arXiv:2403.12239*, 2024.
 - [227] L. Zhang, Z. Liu, Y. Zhou, T. Wu, and J. Sun, “Grounding large language models in real-world environments using imperfect world models,” *International Journal of Advanced Computer Science and Applications*, 2024.
 - [228] U. J. Umoga, E. O. Sodiya, E. D. Ugwuanyi, B. S. Jacks, O. A. Lottu, O. D. Daraojimba, and A. Obaigbena, “Exploring the potential of AI-driven optimization in enhancing network performance and efficiency,” *Magna Scientia Advanced Research and Reviews*, vol. 10, no. 1, pp. 368–378, feb 28 2024.
 - [229] S. K. Mani, Y. Zhou, K. Hsieh, S. Segarra, T. Eberl, E. Azulai, I. Fritzer, R. Chandra, and S. Kandula, “Enhancing network management using code generated by Large Language Models,” in *Proceedings of the 22nd ACM Workshop on Hot Topics in Networks*, 2023, pp. 196–204.
 - [230] D. Wang, S. Wang, X. Jiang, Y. Zhang, Y. Pang, and M. Zhang, “When Large Language Models meet optical networks: Paving the way for automation,” *Electronics*, vol. 13, no. 13, pp. 2529–2529, jun 27 2024.
 - [231] H. Zhou, C. Hu, D. Yuan, Y. Yuan, D. Wu, X. Liu, and C. Zhang, “Large Language Model (LLM)-enabled in-context learning for wireless network optimization: A case study of power control,” *arXiv preprint arXiv:2408.00214*, 2024.
 - [232] K. Patil and B. Desai, “Intelligent network optimization in cloud environments with generative AI and LLMs,” 2024.
 - [233] S. Long, F. Tang, Y. Li, T. Tan, Z. Jin, M. Zhao, and N. Kato, “6G comprehensive intelligence: Network operations and optimization based on Large Language Models,” *arXiv preprint arXiv:2404.18373*, 2024.
 - [234] A. Remadi, K. El Hage, Y. Hobeika, and F. Bugiotti, “To prompt or not to prompt: Navigating the use of Large Language Models for integrating and modeling heterogeneous data,” *Data & Knowledge Engineering*, vol. 152, p. 102313, 2024.
 - [235] M. Hassanin, M. Keshk, S. Salim, M. Alsubaie, and D. Sharma, “PLLM-CS: Pre-trained Large Language Model (LLM) for cyber threat detection in satellite networks,” *Ad Hoc Networks*, vol. 166, p. 103645, 2025.
 - [236] K. B. Kan, H. Mun, G. Cao, and Y. Lee, “Mobile-Llama: Instruction fine-tuning open-source LLM for network analysis in 5G networks,” *IEEE Network*, 2024.
 - [237] J. Zhang, H. Bu, H. Wen, Y. Chen, L. Li, and H. Zhu, “When LLMs meet cybersecurity: A systematic literature review,” *arXiv preprint arXiv:2405.03644*, 2024.
 - [238] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. d. O. Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman, A. Ray, R. Puri, G. Krueger, M. Petrov, H. Khlaaf, G. Sastry, P. Mishkin, B. Chan, S. Gray, N. Ryder, M. Pavlov, A. Power, L. Kaiser, M. Bavarian, C. Winter, P. Tillet, F. P. Such, D. Cummings, M. Plappert, F. Chantzis, E. Barnes, A. Herbert-Voss, W. H. Guss, A. Nichol, A. Paino, N. Tezak, J. Tang, I. Babuschkin, S. Balaji, S. Jain, W. Saunders, C. Hesse, A. N. Carr, J. Leike, J. Achiam, V. Misra, E. Morikawa, A. Radford, M. Knight, M. Brundage, M. Murati, K. Mayer, P. Welinder, B. McGrew, D. Amodei, S. McCandlish, I. Sutskever, and W. Zaremba, “Evaluating Large Language Models trained on code,” *arXiv preprint arXiv:2107.03374*, 2021.
 - [239] T. An, Y. Zhou, H. Zou, and J. Yang, “IoT-LLM: Enhancing real-world IoT task reasoning with Large Language Models,” *arXiv preprint arXiv:2410.02429*, 2024.
 - [240] M. Zong, A. Hekmati, M. Guastalla, Y. Li, and B. Krishnamachari, “Integrating Large Language Models with Internet of Things: Applications,” *Discover Internet of Things*, vol. 5, no. 1, p. 2, 2025.
 - [241] X. Ouyang and M. Srivastava, “LLMSense: Harnessing LLMs for high-level reasoning over spatiotemporal sensor traces,” *arXiv preprint arXiv:2403.19857*, 2024.
 - [242] C. Martin-Turrero, M. Bouvier, M. Breitenstein, P. Zanuttigh, and V. Parret, “ALERT-Transformer: Bridging asynchronous and synchronous machine learning for real-time event-based spatio-temporal data,” *arXiv preprint arXiv:2402.01393*, 2024.
 - [243] H. Xu, P. Zhou, R. Tan, M. Li, and G. Shen, “Limu-Bert: Unleashing the potential of unlabeled data for IMU sensing applications,” in *Proceedings of the 19th ACM Conference on Embedded Networked Sensor Systems*, 2021, pp. 220–233.
 - [244] Q. Huang, Z. Zhou, K. Yang, G. Lin, Z. Yi, and Y. Wang, “LeRet: Language-empowered retentive network for time series forecasting,” in *Proceedings of the International Joint Conference on Artificial Intelligence*, 2024.
 - [245] C. Liu, Q. Xu, H. Miao, S. Yang, L. Zhang, C. Long, Z. Li, and R. Zhao, “TimeCMA: Towards LLM-empowered time series forecast-

- ing via cross-modality alignment,” *arXiv preprint arXiv:2406.01638*, 2024.
- [246] X. Liu, J. Hu, Y. Li, S. Diao, Y. Liang, B. Hooi, and R. Zimmermann, “UniTime: A language-empowered unified model for cross-domain time series forecasting,” in *Proceedings of the ACM on Web Conference 2024*, 2024, pp. 4095–4106.
- [247] D. Cao, F. Jia, S. O. Arik, T. Pfister, Y. Zheng, W. Ye, and Y. Liu, “Tempo: Prompt-based generative pre-trained transformer for time series forecasting,” *arXiv preprint arXiv:2310.04948*, 2023.
- [248] C. Sun, H. Li, Y. Li, and S. Hong, “TEST: Text prototype aligned embedding to activate LLM’s ability for time series,” *arXiv preprint arXiv:2308.08241*, 2023.
- [249] Y.-N. Chuang, S. Li, J. Yuan, G. Wang, K.-H. Lai, L. Yu, S. Ding, C.-Y. Chang, Q. Tan, D. Zha, and X. Hu, “Understanding different design choices in training large time series models,” *arXiv preprint arXiv:2406.14045*, 2024.
- [250] H. Xue and F. D. Salim, “PromptCast: A new prompt-based learning paradigm for time series forecasting,” *IEEE Transactions on Knowledge and Data Engineering*, 2023.
- [251] F. Jia, K. Wang, Y. Zheng, D. Cao, and Y. Liu, “GPT4MTS: Prompt-based Large Language Model for multimodal time-series forecasting,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, no. 21, 2024, pp. 23 343–23 351.
- [252] N. Gruver, M. Finzi, S. Qiu, and A. G. Wilson, “Large Language Models are zero-shot time series forecasters,” *Advances in Neural Information Processing Systems*, vol. 36, 2024.
- [253] G. Civitarese, M. Fiori, P. Choudhary, and C. Bettini, “Large Language Models are zero-shot recognizers for activities of daily living,” *arXiv preprint arXiv:2407.01238*, 2024.
- [254] Z. Yang, L. Li, K. Lin, J. Wang, C.-C. Lin, Z. Liu, and L. Wang, “The dawn of LLMs: Preliminary explorations with FPT-4v (ision),” *arXiv preprint arXiv:2309.17421*, vol. 9, no. 1, p. 1, 2023.
- [255] Microsoft. (2024) Bing image search api. Accessed on 19 July 2024. [Online]. Available: <https://www.microsoft.com/en-us/bing/api/s/bing-image-search-api>
- [256] A. K. Jadoon, C. Yu, and Y. Shi, “ContextMate: A context-aware smart agent for efficient data analysis,” *CCF Transactions on Pervasive Computing and Interaction*, pp. 1–29, 2024.
- [257] H. Zhao, X. Zhang, J. Wei, Y. Xu, Y. He, S. Sun, and C. You, “Timeseriesscientist: A general-purpose ai agent for time series analysis,” *arXiv preprint arXiv:2510.01538*, 2025.
- [258] M. Jin, S. Wang, L. Ma, Z. Chu, J. Y. Zhang, X. Shi, P.-Y. Chen, Y. Liang, Y.-F. Li, S. Pan, and Q. Wen, “Time-LLM: Time series forecasting by reprogramming Large Language Models,” *arXiv preprint arXiv:2310.01728*, 2023.
- [259] C. Chang, W.-Y. Wang, W.-C. Peng, and T.-F. Chen, “Llm4ts: Aligning pre-trained llms as data-efficient time-series forecasters,” *ACM Transactions on Intelligent Systems and Technology*, vol. 16, no. 3, pp. 1–20, 2025.
- [260] K. Post, R. Kuchida, M. Olapade, Z. Yin, P. Nurmi, and H. Flores, “ContextLLM: Meaningful context reasoning from multi-sensor and multi-device data using LLMs,” in *International Workshop on Mobile Computing Systems and Applications*. ACM, 2025.
- [261] Z. Li, L. Xia, J. Tang, Y. Xu, L. Shi, L. Xia, D. Yin, and C. Huang, “UrbanGPT: Spatio-temporal Large Language Models,” *arXiv preprint arXiv:2403.00813*, 2024.
- [262] Y. Bian, X. Ju, J. Li, Z. Xu, D. Cheng, and Q. Xu, “Multi-patch prediction: Adapting LLMs for time series representation learning,” *arXiv preprint arXiv:2402.04852*, 2024.
- [263] Y. Chen, X. Wang, and G. Xu, “GatGPT: A pre-trained Large Language Model with graph attention network for spatiotemporal imputation,” *arXiv preprint arXiv:2311.14332*, 2023.
- [264] A. Narayan, I. Chami, L. Orr, S. Arora, and C. Ré, “Can foundation models wrangle your data?” *arXiv preprint arXiv:2205.09911*, 2022.
- [265] Y. Wang, H. Wu, J. Dong, Y. Liu, M. Long, and J. Wang, “Deep time series models: A comprehensive survey and benchmark,” *arXiv preprint arXiv:2407.13278*, 2024.
- [266] M. E. Barnes, T. A. Ghaleb, and S. Hassan, “Task-aware reduction for scalable llm-database systems,” *arXiv preprint arXiv:2510.11813*, 2025.
- [267] K. Patherya, A. Dhekne, and F. Romero, “Flash-fusion: Enabling expressive, low-latency queries on iot sensor streams with llms,” *arXiv preprint arXiv:2511.11885*, 2025.
- [268] X. Zhang, J. Zhang, B. Rekabdar, Y. Zhou, P. Wang, and K. Liu, “Dynamic and adaptive feature generation with LLM,” *arXiv preprint arXiv:2406.03505*, 2024.
- [269] W. Jiao, W. Wang, J.-t. Huang, X. Wang, and Z. Tu, “Is ChatGPT a good translator? a preliminary study,” *arXiv preprint arXiv:2301.08745*, vol. 1, no. 10, 2023.
- [270] J. Wei, X. Wang, D. Schuurmans, M. Bosma, F. Xia, E. Chi, Q. V. Le, and D. Zhou, “Chain-of-thought prompting elicits reasoning in Large Language Models,” *Advances in neural information processing systems*, vol. 35, pp. 24 824–24 837, 2022.
- [271] Z. Li, Y. Du, J. Hu, X. Wan, and A. Gao, “Self-instructed derived prompt generation meets in-context learning: Unlocking new potential of black-box LLMs,” *arXiv preprint arXiv:2409.01552*, 2024.
- [272] Y. Liu, S. Tao, W. Meng, F. Yao, X. Zhao, and H. Yang, “LogPrompt: Prompt engineering towards zero-shot and interpretable log analysis,” in *Proceedings of the IEEE/ACM International Conference on Software Engineering: Companion Proceedings*, 2024, pp. 364–365.
- [273] J. Heo, V. Piratla, K. Lee, H. Joh, and A. Weller, “LLMs on interactive feature collections with implicit dynamic decision strategy,” in *Proceedings of the International Conference on Computational Linguistics*, 2025, pp. 786–811.
- [274] T. Wang, Z. Zhao, and K. Wu, “Exploiting LLM embeddings for content-based IoT anomaly detection,” in *2024 IEEE Pacific Rim Conference on Communications, Computers and Signal Processing (PACRIM)*. IEEE, 2024, pp. 1–6.
- [275] J. A. Miller, M. Aldosari, F. Saeed, N. H. Barna, S. Rana, I. B. Arpinar, and N. Liu, “A survey of deep learning and foundation models for time series forecasting,” *arXiv preprint arXiv:2401.13912*, 2024.
- [276] T. Tu, Z. He, Z. Zheng, Z. Zheng, J. Jiang, Y. Gong, C. Hu, and D. Cheng, “Towards lifelong unseen task processing with a lightweight unlabeled data schema for AIoT,” *IEEE Internet of Things Journal*, 2024.
- [277] S. Diederich, A. Brendel, S. Morana, and L. M. Kolbe, “On the design of and interaction with conversational agents: An organizing and assessing review of human-computer interaction research,” *Journal of the Association for Information Systems*, vol. 23, no. 1, pp. 96–138, jan 1 2022.
- [278] T. Wu, M. Terry, and C. J. Cai, “Ai Chains: Transparent and Controllable Human-AI Interaction by Chaining Large Language Model Prompts,” in *CHI Conference on Human Factors in Computing Systems*, apr 28 2022, pp. 1–22.
- [279] F. Alwahedi, A. Aldhaheri, M. A. Ferrag, A. Battah, and N. Tihanyi, “Machine learning techniques for IoT security: Current research and future vision with generative AI and Large Language Models,” *Internet of Things and Cyber-Physical Systems*, vol. 4, pp. 167–185, jan 1 2024.
- [280] D. Rodríguez, I. A. Yang, J. M. Del Álamo, and N. Sadeh, “Large Language Models: A new approach for privacy policy analysis at scale,” *Computing*, 2024.
- [281] T. Singh, H. Aditya, V. K. Madiseti, and A. Bahga, “Whispered tuning: Data privacy preservation in fine-tuning LLMs through differential privacy,” *Journal of Software Engineering and Applications*, vol. 17, no. 01, pp. 1–22, 2024.
- [282] F. Neves, R. R. de Souza, J. C. de Sousa, M. Bonfim, and V. C. Garcia, “Data privacy in the Internet of Things based on anonymization: A review,” *Journal of Computer Security*, vol. 31, no. 3, pp. 261–291, 2023.
- [283] G. M. Garrido, J. Sedlmeir, Ö. Uludağ, I. S. Alaoui, A. Luckow, and F. Matthes, “Revealing the landscape of privacy-enhancing technologies in the context of data markets for the IoT: A systematic literature review,” *Journal of Network and Computer Applications*, vol. 207, pp. 103 465–103 465, jul 16 2022.
- [284] X. Zhang, H. Xu, Z. Ba, Z. Wang, Y. Hong, J. Liu, Z. Qin, and K. Ren, “PrivacyAsst: Safeguarding user privacy in tool-using Large Language Model agents,” *IEEE Transactions on Dependable and Secure Computing*, 2024.
- [285] K. Khatawada, J. Hopper, J. Cheatham, A. Joshi, and S. Baidya, “Large language models in the iot ecosystem – a survey on security challenges and applications,” May 2025. [Online]. Available: <https://arxiv.org/abs/2505.17586>
- [286] S. Schiegg, F. Strohmeier, A. Gerl, and H. Kosch, “Individual privacy levels in query-based anonymization,” in *International Conference on Availability, Reliability and Security*, 2024, pp. 1–8.


- [287] R. Staab, M. Vero, M. Balunovic, and M. Vechev, "Large language models are anonymizers," in *ICLR 2024 Workshop on Reliable and Responsible Foundation Models*, 2024.
- [288] G. Dhingra, S. Sood, Z. M. Wase, A. Bahga, and V. K. Madiseti, "Protecting LLMs against privacy attacks while preserving utility," *Journal of Information Security*, vol. 15, no. 4, pp. 448–473, 2024.
- [289] J. Yu, J. Zhou, Y. Ding, L. Zhang, Y. Guo, and H. Sato, "Textual differential privacy for context-aware reasoning with Large Language Model," in *IEEE Annual Computers, Software, and Applications Conference*. IEEE, 2024, pp. 988–997.
- [290] B. Yan, K. Li, M. Xu, Y. Dong, Y. Zhang, Z. Ren, and X. Cheng, "On protecting the data privacy of Large Language Models (LLMs): A survey," *arXiv preprint arXiv:2403.05156*, 2024.
- [291] B. B. Gupta, A. Gaurav, V. Arya, W. Alhalabi, D. Alsaman, and P. Vijayakumar, "Enhancing user prompt confidentiality in Large Language Models through advanced differential encryption," *Computers & Electrical Engineering*, vol. 116, pp. 109 215–109 215, apr 4 2024.
- [292] M. Güven, "A comprehensive review of Large Language Models in cyber security," *International Journal of Computational and Experimental Science and Engineering*, vol. 10, no. 3, sep 30 2024.
- [293] X. Zhao, X. Leng, L. Wang, N. Wang, and Y. Liu, "Efficient anomaly detection in tabular cybersecurity data using Large Language Models," *Scientific Reports*, vol. 15, no. 1, p. 3344, 2025.
- [294] S. Zhou, Z. Zhou, C. Wang, Y. Liang, L. Wang, J. Zhang, J. Zhang, and C. Lv, "A user-centered framework for data privacy protection using Large Language Models and attention mechanisms," *Applied Sciences*, vol. 14, no. 15, pp. 6824–6824, 2024.
- [295] S. Darzi and A. A. Yavuz, "Counter denial of service for next-generation networks within the Artificial Intelligence and post-quantum era," in *IEEE International Conference on Trust, Privacy and Security in Intelligent Systems, and Applications*. IEEE, 2024, pp. 138–147.
- [296] B. Ahmad, S. Thakur, B. Tan, R. Karri, and H. Pearce, "On hardware security bug code fixes by prompting Large Language Models," *IEEE Transactions on Information Forensics and Security*, 2024.
- [297] D. Saha, S. Tarek, K. Yahyaie, S. K. Saha, J. Zhou, M. Tehranipoor, and F. Farahmandi, "LLM for SOC security: A paradigm shift," *IEEE Access*, 2024.
- [298] W. Kuang, B. Qian, Z. Li, D. Chen, D. Gao, X. Pan, Y. Xie, Y. Li, B. Ding, and J. Zhou, "FederatedScope-LLM: A comprehensive package for fine-tuning Large Language Models in federated learning," in *Proceedings of the ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2024, pp. 5260–5271.
- [299] F. Wu, Z. Li, Y. Li, B. Ding, and J. Gao, "FedBiot: LLM local fine-tuning in federated learning without full model," in *Proceedings of the ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2024, pp. 3345–3355.
- [300] X. Zhang, C. Zhang, T. Li, Y. Huang, X. Jia, M. Hu, J. Zhang, Y. Liu, S. Ma, and C. Shen, "Jailguard: A universal detection framework for prompt-based attacks on llm systems," *ACM Transactions on Software Engineering and Methodology*, 2025.
- [301] L. Lu, X. Gu, S. Pang, S. Liang, H. Zhu, X. Zeng, X. Zheng, and Y. Zhou, "E²at: Multimodal jailbreak defense via dynamic joint optimization for multimodal large language models," 2026. [Online]. Available: <https://arxiv.org/abs/2503.04833>
- [302] L. Burgueño, M. Keet, J. Kienzle, J. Michael, and Ö. Babur, "A human behavior exploration approach using LLMs for cyber-physical systems," in *Proceedings of the ACM/IEEE Conference on Model Driven Engineering Languages and Systems*, 2024, pp. 578–586.
- [303] D. Doumanas, G. Bouchouras, A. Soularidis, K. Kotis, and G. Vouros, "From human-to LLM-centered collaborative ontology engineering," *Applied Ontology*, p. 15705838241305067, 2024.
- [304] S. A. Imran, M. N. H. Khan, S. Biswas, and B. Islam, "LLaSA: Large multimodal agent for human activity analysis through wearable sensors," *arXiv preprint arXiv:2406.14498*, 2024.
- [305] C. M. Fang, V. Danry, N. Whitmore, A. Bao, A. Hutchison, C. Pierce, and P. Maes, "PhysioLLM: Supporting personalized health insights with wearables and Large Language Models," *arXiv preprint arXiv:2406.19283*, 2024.
- [306] S. Kou, C. Yang, and M. Gurusamy, "GIA: LLM-enabled generative intent abstraction to enhance adaptability for intent-driven networks," *Transactions on Cognitive Communications and Networking*, 2025.
- [307] N. Stiennon, L. Ouyang, J. Wu, D. M. Ziegler, R. Lowe, C. Voss, A. Radford, D. Amodei, and P. Christiano, "Learning to summarize from human feedback," *arXiv preprint arXiv:2009.01325*, 2022.
- [308] S. Mo, R. Salakhutdinov, L.-P. Morency, and P. P. Liang, "IoT-LM: Large multisensory language models for the Internet of Things," *arXiv preprint arXiv:2407.09801*, 2024.
- [309] J. Chua, Y. Li, S. Yang, C. Wang, and L. Yao, "AI safety in generative AI Large Language Models: A survey," *arXiv preprint arXiv:2407.18369*, 2024.
- [310] D. Rivkin, F. Hogan, A. Feriani, A. Konar, A. Sigal, X. Liu, and G. Dudek, "AIoT smart home via autonomous LLM agents," *IEEE Internet of Things Journal*, 2024.
- [311] X. Luo, D. Liu, F. Dang, and H. Luo, "Integration of LLMs and the physical world: Research and application," in *Proceedings of the ACM Turing Award Celebration Conference-China 2024*, 2024, pp. 1–5.
- [312] B. Liu, Z. Zhang, P. He, Z. Wang, Y. Xiao, R. Ye, Y. Zhou, W.-S. Ku, and B. Hui, "A survey of lottery ticket hypothesis," *arXiv preprint arXiv:2403.04861*, 2024.
- [313] Z. Yin, Y. Yuan, P. Guo, and P. Zhou, "Backdoor attacks on federated learning with lottery ticket hypothesis," *arXiv preprint arXiv:2109.10512*, 2021.
- [314] A. Fan, B. Gokkaya, M. Harman, M. Lyubarskiy, S. Sengupta, S. Yoo, and J. M. Zhang, "Large language models for software engineering: Survey and open problems," in *2023 IEEE/ACM International Conference on Software Engineering: Future of Software Engineering (ICSE-FoSE)*. IEEE, 2023, pp. 31–53.
- [315] F. Wu, X. Liu, H. Wang, X. Wang, and J. Gao, "On the client preference of LLM fine-tuning in federated learning," *arXiv preprint arXiv:2407.03038*, 2024.
- [316] H. Xue, W. Ren, X. Geng, K. Wei, L. Li, Q. Shao, L. Yang, K. Diao, and L. Xie, "Ideal-llm: Integrating dual encoders and language-adapted llm for multilingual speech-to-text," in *National Conference on Man-Machine Speech Communication*. Springer, 2026, pp. 47–58.
- [317] S. Priya, D. Karthika Renuka, and L. Ashok Kumar, "Robust multi-dialect end-to-end asr model jointly with beam search threshold pruning and llm," *SN Computer Science*, vol. 6, no. 4, pp. 1–12, 2025.
- [318] N. E. Ghazal, A. Caubrière, and V. Vielzeuf, "The speech-llm takes it all: A truly fully end-to-end spoken dialogue state tracking approach," *arXiv preprint arXiv:2510.09424*, 2025.
- [319] W. Zhang, S. Naagar, Z. Ye, P. Tang, X. Zhou, J. Liu, and L. Dai, "Bridging modality gap with large speech and language models for end-to-end speech-to-text translation," in *ICASSP 2025-2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2025, pp. 1–5.
- [320] Y. Zhang, Z. Wang, Z. He, J. Li, G. Mai, J. Lin, C. Wei, and W. Yu, "BB-GeoGPT: A framework for learning a Large Language Model for geographic information science," *Information Processing & Management*, vol. 61, no. 5, p. 103808, 2024.
- [321] Y. Hong, H. Zhen, P. Chen, S. Zheng, Y. Du, Z. Chen, and C. Gan, "3D-LLM: Injecting the 3D world into Large Language Models," *Advances in Neural Information Processing Systems*, vol. 36, p. 20482, 2023.
- [322] W. Zhang, S. Zhou, L. Yang, L. Ou, and Z. Xiao, "WiFiMap+: high-level indoor semantic inference with WiFi human activity and environment," *IEEE Transactions on Vehicular Technology*, vol. 68, no. 8, pp. 7890–7903, 2019.
- [323] P. Shaw, M. Mikusz, P. Nurmi, and N. Davies, "IoT maps: Charting the Internet of Things," in *Proceedings of the 20th International Workshop on Mobile Computing Systems and Applications*, 2019, pp. 105–110.
- [324] Y. Gu, T. Liu, J. Li, F. Ren, Z. Liu, X. Wang, and P. Li, "EmoSense: Data-driven emotion sensing via off-the-shelf WiFi devices," in *IEEE International Conference on Communications*. IEEE, 2018, pp. 1–6.
- [325] H. Niu, H. Li, S. Gao, Y. Li, X. Wei, Y. Chen, W. Yue, W. Zhou, and G. Shen, "Perception-to-cognition tactile sensing based on artificial-intelligence-motivated human full-skin bionic electronic skin," *Advanced Materials*, vol. 34, no. 31, p. 2202622, 2022.
- [326] M. McManus, Y. Cui, J. Z. Zhang, J. Hu, S. K. Moorthy, N. Masronar, E. S. Bentley, M. Medley, and Z. Guan, "Digital twin-enabled domain adaptation for zero-touch UAV networks: Survey and challenges," *Computer Networks*, vol. 236, p. 110000, 2023.

- [327] C. Bettini, O. Brdiczka, K. Henriksen, J. Indulska, D. Nicklas, A. Ranganathan, and D. Riboni, "A survey of context modelling and reasoning techniques," *Pervasive and Mobile Computing*, vol. 6, no. 2, pp. 161–180, 2010.
- [328] S. Salehi, M. Erol-Kantarci, and D. Niyato, "LLM-enabled data transmission in end-to-end semantic communication," *arXiv preprint arXiv:2504.07431*, 2025.
- [329] S. Guo, Y. Wang, J. Ye, A. Zhang, P. Zhang, and K. Xu, "Semantic importance-aware communications with semantic correction using large language models," *IEEE Transactions on Machine Learning in Communications and Networking*, 2025.
- [330] Y. Hong, J. Wu, and R. Morello, "LLM-Twin: mini-giant model-driven beyond 5G digital twin networking framework with semantic secure communication and computation," *Scientific Reports*, vol. 14, no. 1, p. 19065, 2024.
- [331] J. Liu, M. Liu, J.-P. Liu, Z. Ye, Y. Wang, Y. Alexeev, J. Eisert, and L. Jiang, "Towards provably efficient quantum algorithms for large-scale machine-learning models," *Nature Communications*, vol. 15, no. 1, p. 434, 2024.
- [332] B. Aizpurua, S. S. Jahromi, S. Singh, and R. Orus, "Quantum Large Language Models via tensor network disentanglers," *arXiv preprint arXiv:2410.17397*, 2024.
- [333] C. Y. Kim, C. P. Lee, and B. Mutlu, "Understanding large-language model (LLM)-powered human-robot interaction," in *Proceedings of the 2024 ACM/IEEE International Conference on Human-Robot Interaction*, 2024, pp. 371–380.
- [334] H. Liu, Y. Zhu, K. Kato, A. Tsukahara, I. Kondo, T. Aoyama, and Y. Hasegawa, "Enhancing the LLM-based robot manipulation through human-robot collaboration," *IEEE Robotics and Automation Letters*, 2024.
- [335] X. Yuan, L. Ding, L. Zhang, X. Li, and D. O. Wu, "ES Attack: Model stealing against deep neural networks without data hurdles," *IEEE Transactions on Emerging Topics in Computational Intelligence*, vol. 6, no. 5, pp. 1258–1270, 2022.
- [336] J. Zhu, H. Yin, P. Deng, A. Almeida, and S. Zhou, "Confidential computing on nvidia hopper gpus: A performance benchmark study," *arXiv preprint arXiv:2409.03992*, 2024.
- [337] D. Li, Z. Zhang, M. Yao, Y. Cai, Y. Guo, and X. Chen, "TEESlice: Protecting sensitive neural network models in trusted execution environments when attackers have pre-trained models," *ACM Transactions on Software Engineering and Methodology*, Dec. 2024. [Online]. Available: <https://doi.org/10.1145/3707453>
- [338] H. Yang, D. Zhang, Y. Zhao, Y. Li, and Y. Liu, "A first look at efficient and secure on-device LLM inference against KV leakage," in *Proceedings of the 19th Workshop on Mobility in the Evolving Internet Architecture*, ser. MobiArch '24. New York, NY, USA: Association for Computing Machinery, 2024, p. 13–18. [Online]. Available: <https://doi.org/10.1145/3691555.3696827>
- [339] M. Cuadros, A. Álvarez, N. Pérez, J. M. Martín, P. Turón, E. Zotova, H. Arzelus, J. Arellano, and A. del Pozo, "Emphasis: Empowering decision making with higher productivity by means of hyperautomation," in *Proceedings of the Seminar of the Spanish Society for Natural Language Processing (SEPLN-CEDI-PD)*. CEUR, 2024.
- [340] Z. Zheng, S. Xie, H.-N. Dai, W. Chen, X. Chen, J. Weng, and M. Imran, "An overview on smart contracts: Challenges, advances and platforms," *Future Generation Computer Systems*, vol. 105, pp. 475–491, 2020.
- [341] Y. Hu, Y. Koren, and C. Volinsky, "Collaborative filtering for implicit feedback datasets," in *IEEE International Conference on Data Mining*. IEEE, 2008, pp. 263–272.
- [342] M. Slater, "Implicit learning through embodiment in immersive virtual reality," in *Virtual, augmented, and mixed realities in education*. Springer, 2017, pp. 19–33.
- [343] B. C. Das, M. H. Amini, and Y. Wu, "Security and privacy challenges of large language models: A survey," *ACM Computing Surveys*, vol. 57, no. 6, pp. 1–39, 2025.
- [344] D. Adhikari, W. Jiang, J. Zhan, Z. He, D. B. Rawat, U. Aickelin, and H. A. Khorshidi, "A comprehensive survey on imputation of missing data in Internet of Things," *ACM Computing Surveys*, vol. 55, no. 7, pp. 1–38, 2022.
- [345] A. Jobin, M. Ienca, and E. Vayena, "The global landscape of AI ethics guidelines," *Nature Machine Intelligence*, vol. 1, no. 9, p. 389, 2019.

Biographies



Fatemeh Sarhaddi received the Ph.D degree from University of Turku, Finland, in 2024. She is currently a Postdoctoral Researcher at the Pervasive Data Science research group, Department of Computer Science, University of Helsinki. Her research interests include Internet of Things, digital health, and federated learning. Email: fatemeh.sarhaddi@helsinki.fi



Ngoc Thi Nguyen received the Ph.D degree from Singapore University of Technology and Design, Singapore, in 2021. She is currently a Postdoctoral Researcher at the Pervasive Data Science research group, Department of Computer Science, University of Helsinki. Her research interests include pervasive data science, Artificial Intelligence of Things (AIoT), and smart multi-device systems. Email: ngoc.t.nguyen@helsinki.fi



Agustín Zuniga received the Ph.D degree from University of Helsinki, Finland, in 2023. He is currently a Postdoctoral Researcher at the Pervasive Data Science research group, Department of Computer Science, University of Helsinki. His doctoral thesis received the Outstanding Thesis award, University of Helsinki. His research interests include pervasive data science, Internet of Things, and Artificial Intelligence of Things (AIoT). Email: agustin.zuniga@helsinki.fi



Musfira Khan is a doctoral researcher at the Pervasive Data Science research group, Department of Computer Science, University of Helsinki. Her research interests include artificial intelligence, distributed systems, Internet of Things (IoT), and pervasive data science. Email: musfira.khan@helsinki.fi



is an ACM Distinguished Scientist and a member of Academia Europaea. Email: panhui@ust.hk



Sasu Tarkoma (Senior Member, IEEE) received the Ph.D. degree in computer science from the University of Helsinki, Finland, in 2006. He is currently the Dean of the Faculty of Science and a Full Professor at the Department of Computer Science, University of Helsinki. He is also Chairman of the Scientific Advisory Board for Finnish Defence (MATINE) and the Director of the Helsinki Center for Data Science (HiDATA). His research has received several Best Paper awards and mentions, like at IEEE PerCom, ACM CCR and ACM SenSys. His research interests include distributed systems, data

analytics, mobile and ubiquitous computing, artificial intelligence, and 6G.
Email: [sasutarkoma@helsinki.fi](mailto:sasu.tarkoma@helsinki.fi)



Huber Flores (Member, IEEE) received the Ph.D degree from University of Helsinki, Finland, in 2015. He is currently an Associate Professor with the Institute of Computer Science, University of Tartu and a Docent (Affiliate Professor) at the University of Helsinki, Finland. His research interests include mobile and pervasive computing, distributed systems, and mobile cloud computing. Email: huber.flores@ut.ee



Petteri Nurmi received the Ph.D degree from University of Helsinki, Finland, in 2009. He is currently a Full Professor at the Department of Computer Science, University of Helsinki, and a former Lecturer at Lancaster University, UK. His research has received several the Best Paper award at IEEE PerCom and the Test of Time award at ACM SenSys. His research interests include distributed systems, Artificial Intelligence of Things (AIoT), and pervasive data science. Email: petteri.nurmi@helsinki.fi